

```
import json

import logging

import time

import uuid

from datetime import datetime

from typing import Any, Dict, List, Optional


import yaml

from pydantic import BaseModel

from swarm_models.tiktoken_wrapper import TikTokenizer


logger = logging.getLogger(__name__)


class MemoryMetadata(BaseModel):

    """Metadata for memory entries"""

    timestamp: Optional[float] = time.time()

    role: Optional[str] = None

    agent_name: Optional[str] = None

    session_id: Optional[str] = None

    memory_type: Optional[str] = None # 'short_term' or 'long_term'

    token_count: Optional[int] = None

    message_id: Optional[str] = str(uuid.uuid4())
```

```
class MemoryEntry(BaseModel):
```

```
    """Single memory entry with content and metadata"""
```

```
    content: Optional[str] = None
```

```
    metadata: Optional[MemoryMetadata] = None
```

```
class MemoryConfig(BaseModel):
```

```
    """Configuration for memory manager"""
```

```
    max_short_term_tokens: Optional[int] = 4096
```

```
    max_entries: Optional[int] = None
```

```
    system_messages_token_buffer: Optional[int] = 1000
```

```
    enable_long_term_memory: Optional[bool] = False
```

```
    auto_archive: Optional[bool] = True
```

```
    archive_threshold: Optional[float] = 0.8 # Archive when 80% full
```

```
class MemoryManager:
```

```
    """
```

```
    Manages both short-term and long-term memory for an agent, handling token limits,  
    archival, and context retrieval.
```

```
    Args:
```

```
        config (MemoryConfig): Configuration for memory management
```

```
        tokenizer (Optional[Any]): Tokenizer to use for token counting
```

long_term_memory (Optional[Any]): Vector store or database for long-term storage

"""

```
def __init__(
    self,
    config: MemoryConfig,
    tokenizer: Optional[Any] = None,
    long_term_memory: Optional[Any] = None,
):
    self.config = config

    self.tokenizer = tokenizer or TikTokenizer()

    self.long_term_memory = long_term_memory

    # Initialize memories

    self.short_term_memory: List[MemoryEntry] = []
    self.system_messages: List[MemoryEntry] = []

    # Memory statistics

    self.total_tokens_processed: int = 0
    self.archived_entries_count: int = 0

def create_memory_entry(
    self,
    content: str,
    role: str,
    agent_name: str,
```

```

    session_id: str,

    memory_type: str = "short_term",

) -> MemoryEntry:

    """Create a new memory entry with metadata"""

    metadata = MemoryMetadata(

        timestamp=time.time(),

        role=role,

        agent_name=agent_name,

        session_id=session_id,

        memory_type=memory_type,

        token_count=self.tokenizer.count_tokens(content),

    )

    return MemoryEntry(content=content, metadata=metadata)

```

```

def add_memory(

    self,

    content: str,

    role: str,

    agent_name: str,

    session_id: str,

    is_system: bool = False,

```

```

) -> None:

    """Add a new memory entry to appropriate storage"""

    entry = self.create_memory_entry(

        content=content,

        role=role,

```

```
agent_name=agent_name,  
session_id=session_id,  
memory_type="system" if is_system else "short_term",  
)
```

```
if is_system:
```

```
    self.system_messages.append(entry)
```

```
else:
```

```
    self.short_term_memory.append(entry)
```

```
# Check if archiving is needed
```

```
if self.should_archive():
```

```
    self.archive_old_memories()
```

```
self.total_tokens_processed += entry.metadata.token_count
```

```
def get_current_token_count(self) -> int:
```

```
    """Get total tokens in short-term memory"""
```

```
    return sum(  
        entry.metadata.token_count  
        for entry in self.short_term_memory  
    )
```

```
def get_system_messages_token_count(self) -> int:
```

```
    """Get total tokens in system messages"""
```

```
    return sum(  
        entry.metadata.token_count  
        for entry in self.system_messages  
    )
```

```
        entry.metadata.token_count
    for entry in self.system_messages
)
```

```
def should_archive(self) -> bool:
```

```
    """Check if archiving is needed based on configuration"""
```

```
    if not self.config.auto_archive:
```

```
        return False
```

```
    current_usage = (
```

```
        self.get_current_token_count()
```

```
        / self.config.max_short_term_tokens
```

```
)
```

```
    return current_usage >= self.config.archive_threshold
```

```
def archive_old_memories(self) -> None:
```

```
    """Move older memories to long-term storage"""
```

```
    if not self.long_term_memory:
```

```
        logger.warning(
```

```
            "No long-term memory storage configured for archiving"
```

```
)
```

```
    return
```

```
while self.should_archive():
```

```
    # Get oldest non-system message
```

```
    if not self.short_term_memory:
```

break

oldest_entry = self.short_term_memory.pop(0)

Store in long-term memory

self.store_in_long_term_memory(oldest_entry)

self.archived_entries_count += 1

def store_in_long_term_memory(self, entry: MemoryEntry) -> None:

"""Store a memory entry in long-term memory"""

if self.long_term_memory is None:

logger.warning(

"Attempted to store in non-existent long-term memory"

)

return

try:

self.long_term_memory.add(str(entry.model_dump()))

except Exception as e:

logger.error(f"Error storing in long-term memory: {e}")

Re-add to short-term if storage fails

self.short_term_memory.insert(0, entry)

def get_relevant_context(

self, query: str, max_tokens: Optional[int] = None

) -> str:

```
"""
```

Get relevant context from both memory types

Args:

query (str): Query to match against memories

max_tokens (Optional[int]): Maximum tokens to return

Returns:

str: Combined relevant context

```
"""
```

```
contexts = []
```

```
# Add system messages first
```

```
for entry in self.system_messages:
```

```
    contexts.append(entry.content)
```

```
# Add short-term memory
```

```
for entry in reversed(self.short_term_memory):
```

```
    contexts.append(entry.content)
```

```
# Query long-term memory if available
```

```
if self.long_term_memory is not None:
```

```
    long_term_context = self.long_term_memory.query(query)
```

```
    if long_term_context:
```

```
        contexts.append(str(long_term_context))
```



```
# Combine and truncate if needed
```

```
combined = "\n".join(contexts)
```

```
if max_tokens:
```

```
    combined = self.truncate_to_token_limit(
```

```
        combined, max_tokens
```

```
    )
```

```
return combined
```

```
def truncate_to_token_limit(
```

```
    self, text: str, max_tokens: int
```

```
) -> str:
```

```
    """Truncate text to fit within token limit"""
```

```
    current_tokens = self.tokenizer.count_tokens(text)
```

```
    if current_tokens <= max_tokens:
```

```
        return text
```

```
# Truncate by splitting into sentences and rebuilding
```

```
sentences = text.split(". ")
```

```
result = []
```

```
current_count = 0
```

```
for sentence in sentences:
```

```
    sentence_tokens = self.tokenizer.count_tokens(sentence)
```

```
    if current_count + sentence_tokens <= max_tokens:
```

```
result.append(sentence)
```

```
current_count += sentence_tokens
```

```
else:
```

```
break
```

```
return ". ".join(result)
```

```
def clear_short_term_memory(
```

```
    self, preserve_system: bool = True
```

```
) -> None:
```

```
    """Clear short-term memory with option to preserve system messages"""
```

```
    if not preserve_system:
```

```
        self.system_messages.clear()
```

```
    self.short_term_memory.clear()
```

```
    logger.info(
```

```
        "Cleared short-term memory"
```

```
        + " (preserved system messages)"
```

```
        if preserve_system
```

```
        else ""
```

```
)
```

```
def get_memory_stats(self) -> Dict[str, Any]:
```

```
    """Get detailed memory statistics"""
```

```
    return {
```

```
        "short_term_messages": len(self.short_term_memory),
```

```
        "system_messages": len(self.system_messages),
```

```

"current_tokens": self.get_current_token_count(),
"system_tokens": self.get_system_messages_token_count(),
"max_tokens": self.config.max_short_term_tokens,
"token_usage_percent": round(
    (
        self.get_current_token_count()
        / self.config.max_short_term_tokens
    )
    * 100,
    2,
),
"has_long_term_memory": self.long_term_memory is not None,
"archived_entries": self.archived_entries_count,
"total_tokens_processed": self.total_tokens_processed,
}

```

```

def save_memory_snapshot(self, file_path: str) -> None:

```

```

    """Save current memory state to file"""

```

```

    try:

```

```

        data = {

```

```

            "timestamp": datetime.now().isoformat(),

```

```

            "config": self.config.model_dump(),

```

```

            "system_messages": [

```

```

                entry.model_dump()

```

```

                for entry in self.system_messages

```

```

            ],

```

```

"short_term_memory": [
    entry.model_dump()
    for entry in self.short_term_memory
],
"stats": self.get_memory_stats(),
}

```

```

with open(file_path, "w") as f:
    if file_path.endswith(".yaml"):
        yaml.dump(data, f)
    else:
        json.dump(data, f, indent=2)

```

```

logger.info(f"Saved memory snapshot to {file_path}")

```

```

except Exception as e:
    logger.error(f"Error saving memory snapshot: {e}")
    raise

```

```

def load_memory_snapshot(self, file_path: str) -> None:

```

```

    """Load memory state from file"""

```

```

    try:

```

```

        with open(file_path, "r") as f:
            if file_path.endswith(".yaml"):
                data = yaml.safe_load(f)
            else:

```

```
data = json.load(f)
```

```
self.config = MemoryConfig(**data["config"])
```

```
self.system_messages = [
```

```
    MemoryEntry(**entry)
```

```
    for entry in data["system_messages"]
```

```
]
```

```
self.short_term_memory = [
```

```
    MemoryEntry(**entry)
```

```
    for entry in data["short_term_memory"]
```

```
]
```

```
logger.info(f"Loaded memory snapshot from {file_path}")
```

```
except Exception as e:
```

```
    logger.error(f"Error loading memory snapshot: {e}")
```

```
    raise
```

```
def search_memories(
```

```
    self, query: str, memory_type: str = "all"
```

```
) -> List[MemoryEntry]:
```

```
    """
```

```
    Search through memories of specified type
```

```
    Args:
```

```
        query (str): Search query
```

memory_type (str): Type of memories to search ("short_term", "system", "long_term", or "all")

Returns:

List[MemoryEntry]: Matching memory entries

"""

results = []

if memory_type in ["short_term", "all"]:

results.extend(

[

entry

for entry in self.short_term_memory

if query.lower() in entry.content.lower()

]

)

if memory_type in ["system", "all"]:

results.extend(

[

entry

for entry in self.system_messages

if query.lower() in entry.content.lower()

]

)

if (

```

memory_type in ["long_term", "all"]

and self.long_term_memory is not None

):

long_term_results = self.long_term_memory.query(query)

if long_term_results:

    # Convert long-term results to MemoryEntry format

    for result in long_term_results:

        content = str(result)

        metadata = MemoryMetadata(

            timestamp=time.time(),

            role="long_term",

            agent_name="system",

            session_id="long_term",

            memory_type="long_term",

            token_count=self.tokenizer.count_tokens(

                content

            ),

        )

        results.append(

            MemoryEntry(

                content=content, metadata=metadata

            )

        )

return results

```

```

def get_memory_by_timeframe(
    self, start_time: float, end_time: float
) -> List[MemoryEntry]:
    """Get memories within a specific timeframe"""
    return [
        entry
        for entry in self.short_term_memory
        if start_time <= entry.metadata.timestamp <= end_time
    ]

```

```

def export_memories(
    self, file_path: str, format: str = "json"
) -> None:
    """Export memories to file in specified format"""
    data = {
        "system_messages": [
            entry.model_dump() for entry in self.system_messages
        ],
        "short_term_memory": [
            entry.model_dump() for entry in self.short_term_memory
        ],
        "stats": self.get_memory_stats(),
    }

```

```

with open(file_path, "w") as f:

```

```

    if format == "yaml":

```



```
yaml.dump(data, f)
```

```
else:
```

```
    json.dump(data, f, indent=2)
```