

```
import json

import sqlite3

from typing import List, Optional

from contextlib import contextmanager

from datetime import datetime

import structlog

from fastapi import FastAPI, HTTPException, Request

from fastapi.middleware.cors import CORSMiddleware

from fastapi.middleware.gzip import GZipMiddleware

from fastapi.middleware.trustedhost import TrustedHostMiddleware

from pydantic import BaseModel

from opentelemetry import trace, metrics

from opentelemetry.exporter.otlp.proto.grpc.trace_exporter import (
    OTLPSpanExporter,
)

from opentelemetry.sdk.trace import TracerProvider

from opentelemetry.sdk.trace.export import BatchSpanProcessor

from opentelemetry.sdk.metrics import MeterProvider

from opentelemetry.sdk.metrics.export import (
    PeriodicExportingMetricReader,
)

from opentelemetry.exporter.otlp.proto.grpc.metric_exporter import (
    OTLPMetricExporter,
)

from opentelemetry.instrumentation.fastapi import FastAPIInstrumentor
```

```
# from opentelemetry.instrumentation.sqlite3 import SQLite3Instrumentor

from prometheus_client import Counter, Histogram

import uvicorn

from mcs.main import MedicalCoderSwarm


# Configure structured logging

logger = structlog.get_logger()


# Configure OpenTelemetry

tracer_provider = TracerProvider()

otlp_span_exporter = OTLPSpanExporter()

span_processor = BatchSpanProcessor(otlp_span_exporter)

tracer_provider.add_span_processor(span_processor)

trace.set_tracer_provider(tracer_provider)


# Configure metrics

metric_reader = PeriodicExportingMetricReader(OTLPMetricExporter())

meter_provider = MeterProvider(metric_readers=[metric_reader])

metrics.set_meter_provider(meter_provider)


meter = metrics.get_meter(__name__)

request_counter = meter.create_counter(

    name="api_requests_total",

    description="Total number of API requests",

    unit="1",

)
```

```
# Initialize Prometheus metrics

REQUEST_TIME = Histogram(
    "request_processing_seconds",
    "Time spent processing request",
    ["endpoint"],
)

ERROR_COUNTER = Counter(
    "api_errors_total",
    "Total number of API errors",
    ["endpoint", "error_type"],
)
```

```
# Database configuration
```

```
DB_POOL_SIZE = 5
```

```
db_path = "medical_coder.db"
```

```
class DatabasePool:
```

```
    def __init__(self, database_path: str, pool_size: int):
```

```
        self.database_path = database_path
```

```
        self.pool_size = pool_size
```

```
        self.connections = []
```

```
        self.initialize_pool()
```

```
    def initialize_pool(self):
```

```
for _ in range(self.pool_size):  
    conn = sqlite3.connect(self.database_path)  
    conn.row_factory = sqlite3.Row  
    self.connections.append(conn)
```

```
@contextmanager
```

```
def get_connection(self):  
    if not self.connections:  
        conn = sqlite3.connect(self.database_path)  
        conn.row_factory = sqlite3.Row  
    else:  
        conn = self.connections.pop()  
  
    try:  
        yield conn  
    finally:  
        self.connections.append(conn)
```

```
db_pool = DatabasePool(db_path, DB_POOL_SIZE)
```

```
# Initialize FastAPI app with additional configuration
```

```
app = FastAPI(  
    title="MedicalCoderSwarm API",  
    version="1.0.0",  
    docs_url="/api/docs",
```

```
    redoc_url="/api/redoc",
)

# Add middleware

app.add_middleware(
    CORSMiddleware,
    allow_origins=["*"],
    allow_credentials=True,
    allow_methods=["*"],
    allow_headers=["*"],
)

app.add_middleware(GZipMiddleware, minimum_size=1000)

app.add_middleware(
    TrustedHostMiddleware,
    allowed_hosts=[
        "localhost",
        "127.0.0.1",
    ], # Configure for production
)

# Instrument FastAPI with OpenTelemetry

FastAPIInstrumentor.instrument_app(app)

# SQLite3Instrumentor().instrument()

# Pydantic models
```

```
class PatientCase(BaseModel):  
  
    patient_id: Optional[str] = None  
  
    case_description: Optional[str] = None
```

```
class QueryResponse(BaseModel):  
  
    patient_id: Optional[str] = None  
  
    case_data: Optional[str] = None  
  
    timestamp: datetime = datetime.utcnow()
```

```
class QueryAllResponse(BaseModel):  
  
    patients: Optional[List[QueryResponse]] = None  
  
    total_count: int  
  
    timestamp: datetime = datetime.utcnow()
```

```
class BatchPatientCase(BaseModel):  
  
    cases: Optional[List[PatientCase]] = None
```

```
# Middleware for request tracking
```

```
@app.middleware("http")
```

```
async def add_process_time_header(request: Request, call_next):
```

```
    start_time = datetime.utcnow()
```

```
    response = await call_next(request)
```

```
process_time = (datetime.utcnow() - start_time).total_seconds()
```

```
REQUEST_TIME.labels(endpoint=request.url.path).observe(
```

```
    process_time
```

```
)
```

```
response.headers["X-Process-Time"] = str(process_time)
```

```
return response
```

```
# Enhanced database functions
```

```
def fetch_patient_data(patient_id: str) -> Optional[dict]:
```

```
    with db_pool.get_connection() as conn:
```

```
        try:
```

```
            cursor = conn.cursor()
```

```
            cursor.execute(
```

```
                "SELECT patient_data FROM patients WHERE patient_id = ?",
```

```
                (patient_id,),
```

```
            )
```

```
            row = cursor.fetchone()
```

```
            return json.loads(row[0]) if row else None
```

```
        except sqlite3.Error as e:
```

```
            logger.error(
```

```
                "database_error", error=str(e), patient_id=patient_id
```

```
            )
```

```
            raise HTTPException(
```

```
status_code=500, detail=f"Database error: {str(e)}"
```

```
)
```

```
def save_patient_data(patient_id: str, patient_data: str):
```

```
    with db_pool.get_connection() as conn:
```

```
        try:
```

```
            cursor = conn.cursor()
```

```
            cursor.execute(  
                """  
                INSERT OR REPLACE INTO patients  
                (patient_id, patient_data, created_at, updated_at)  
                VALUES (?, ?, datetime('now'), datetime('now'))  
                """,  
                (patient_id, patient_data),  
            )  
            conn.commit()
```

```
        except sqlite3.Error as e:
```

```
            logger.error(  
                "database_error", error=str(e), patient_id=patient_id  
            )  
            raise HTTPException(  
                status_code=500, detail=f"Database error: {str(e)}"  
            )
```



```
# Enhanced API endpoints
```

```
@app.post("/v1/medical-coder/run", response_model=QueryResponse)
```

```
async def run_medical_coder(
```

```
    patient_case: PatientCase, request: Request
```

```
):
```

```
    tracer = trace.get_tracer(__name__)
```

```
    with tracer.start_as_current_span("run_medical_coder") as span:
```

```
        try:
```

```
            span.set_attribute("patient_id", patient_case.patient_id)
```

```
            logger.info(
```

```
                "processing_patient_case",
```

```
                patient_id=patient_case.patient_id,
```

```
                request_id=request.headers.get("X-Request-ID"),
```

```
            )
```

```
            swarm = MedicalCoderSwarm(
```

```
                patient_id=patient_case.patient_id,
```

```
                max_loops=1,
```

```
                patient_documentation="",
```

```
            )
```

```
            swarm.run(task=patient_case.case_description)
```

```
            swarm_output = swarm.to_dict()
```

```
            save_patient_data(
```

```
                patient_case.patient_id, json.dumps(swarm_output)
```

```
            )
```

```
request_counter.add(1, {"endpoint": "run_medical_coder"})
```

```
return QueryResponse(  
    patient_id=patient_case.patient_id,  
    case_data=json.dumps(swarm_output),  
    timestamp=datetime.utcnow(),  
)
```

```
except Exception as error:
```

```
    ERROR_COUNTER.labels(  
        endpoint="run_medical_coder",  
        error_type=type(error).__name__,  
    ).inc()
```

```
    logger.error(  
        "medical_coder_error",  
        error=str(error),  
        patient_id=patient_case.patient_id,  
    )
```

```
    raise HTTPException(  
        status_code=500,  
        detail=f"Processing error: {str(error)}",  
    )
```

```
@app.get(  
    "/v1/medical-coder/patient/{patient_id}",
```

```

        response_model=QueryResponse,
    )

    async def get_patient_data(patient_id: str, request: Request):

        tracer = trace.get_tracer(__name__)

        with tracer.start_as_current_span("get_patient_data") as span:

            try:

                span.set_attribute("patient_id", patient_id)

                patient_data = fetch_patient_data(patient_id)

                if not patient_data:

                    raise HTTPException(

                        status_code=404, detail="Patient not found"

                    )

                request_counter.add(1, {"endpoint": "get_patient_data"})

                return QueryResponse(

                    patient_id=patient_id,

                    case_data=json.dumps(patient_data),

                    timestamp=datetime.utcnow(),

                )

            except Exception as error:

                ERROR_COUNTER.labels(

                    endpoint="get_patient_data",

                    error_type=type(error).__name__,

                ).inc()

```

```
logger.error(  
    "fetch_patient_error",  
    error=str(error),  
    patient_id=patient_id,  
)  
  
raise
```

```
@app.get("/v1/medical-coder/health")
```

```
async def health_check():
```

```
    """Health check endpoint for monitoring"""
```

```
    try:
```

```
        with db_pool.get_connection() as conn:
```

```
            cursor = conn.cursor()
```

```
            cursor.execute("SELECT 1")
```

```
            return {
```

```
                "status": "healthy",
```

```
                "timestamp": datetime.utcnow(),
```

```
            }
```

```
    except Exception as e:
```

```
        logger.error("health_check_failed", error=str(e))
```

```
        raise HTTPException(  
            status_code=503, detail="Service Unavailable"
```

```
)
```

```
if __name__ == "__main__":

    try:

        uvicorn.run(

            app,

            host="0.0.0.0",

            port=8000,

            workers=4,

            log_config={

                "version": 1,

                "disable_existing_loggers": False,

                "formatters": {

                    "json": {

                        "()": structlog.stdlib.ProcessorFormatter,

                        "processor": structlog.processors.JSONRenderer(),

                    }

                },

                "handlers": {

                    "default": {

                        "class": "logging.StreamHandler",

                        "formatter": "json",

                    }

                },

                "loggers": {

                    "": {

                        "handlers": ["default"],

                        "level": "INFO",
```

```
    }  
    },  
    },  
)
```

```
except Exception as e:
```

```
    logger.error("startup_error", error=str(e))
```

```
    raise
```