

```
import os

import sys

from dataclasses import dataclass

from typing import Any, Callable, Dict, List, Optional, Tuple

from functools import wraps

import inspect

import numpy as np

import psutil

import torch

from loguru import logger

from tenacity import (

    retry,

    retry_if_exception_type,

    stop_after_attempt,

    wait_fixed,

)

import time


# Configure Loguru logger

LOG_LEVEL = os.getenv("LOG_LEVEL", "INFO")

LOG_FORMAT = (

    "<green>{time:YYYY-MM-DD HH:mm:ss.SSS}</green> | "

    "<level>{level: <8}</level> | "

    "<cyan>{name}</cyan>:<cyan>{function}</cyan>:<cyan>{line}</cyan> | "

    "<level>{message}</level>"

)
```

```
# Remove default logger and add custom configuration
```

```
logger.remove()
```

```
logger.add(
```

```
    sys.stderr,
```

```
    format=LOG_FORMAT,
```

```
    level=LOG_LEVEL,
```

```
    backtrace=True,
```

```
    diagnose=True,
```

```
    enqueue=True,
```

```
)
```

```
logger.add(
```

```
    "resource_manager_{time}.log",
```

```
    format=LOG_FORMAT,
```

```
    level=LOG_LEVEL,
```

```
    rotation="100 MB",
```

```
    retention="30 days",
```

```
    compression="zip",
```

```
)
```

```
@dataclass
```

```
class ResourceRequirements:
```

```
    """Data class to store predicted resource requirements."""
```

```
    cpu_memory_bytes: int
```

```
gpu_memory_bytes: int  
requires_gpu: bool  
estimated_runtime_seconds: float
```

```
@dataclass
```

```
class ResourceAllocation:
```

```
    """Data class to store actual resource allocation."""
```

```
    cpu_memory_reserved: int
```

```
    gpu_memory_reserved: int
```

```
    gpu_id: Optional[int]
```

```
    cpu_cores: List[int]
```

```
    allocation_time: float = time.time()
```

```
class InsufficientResourcesError(Exception):
```

```
    """Raised when there are not enough resources to execute the function."""
```

```
    pass
```

```
class MemoryPredictor:
```

```
    """Predicts memory requirements for function execution."""
```

```
    def __init__(self):
```

```
self.history: Dict[str, List[ResourceRequirements]] = {}
```

```
self.sample_size = 3
```

```
logger.info(
```

```
    "Initialized MemoryPredictor with sample size {}",
```

```
    self.sample_size,
```

```
)
```

```
def update_history(
```

```
    self, func_name: str, requirements: ResourceRequirements
```

```
) -> None:
```

```
    """
```

```
    Updates the execution history for a function with new resource requirements.
```

```
    Args:
```

```
        func_name: Name of the function
```

```
        requirements: Resource requirements from the latest execution
```

```
    """
```

```
    try:
```

```
        with logger.contextualize(function=func_name):
```

```
            if func_name not in self.history:
```

```
                self.history[func_name] = []
```

```
                self.history[func_name].append(requirements)
```

```
                # Keep only recent history
```

```
                if len(self.history[func_name]) > self.sample_size:
```

```

self.history[func_name] = self.history[func_name][
    -self.sample_size :
]

logger.debug(
    "Updated history for function {}. Total records: {}".format(
        func_name,
        len(self.history[func_name]),
    )
)

# Calculate and log average requirements
avg_cpu = sum(
    r.cpu_memory_bytes
    for r in self.history[func_name]
) / len(self.history[func_name])
avg_gpu = sum(
    r.gpu_memory_bytes
    for r in self.history[func_name]
) / len(self.history[func_name])

logger.debug(
    "Average requirements - CPU: {} bytes, GPU: {} bytes".format(
        int(avg_cpu),
        int(avg_gpu),
    )
)

```

except Exception as e:

```
    logger.exception(
        "Error updating history for function {}: {}".format(
            func_name,
            str(e),
        )
    )
```

def get_historical_requirements(

self, func_name: str

) -> Optional[ResourceRequirements]:

"""

Gets average historical resource requirements for a function.

Args:

func_name: Name of the function

Returns:

Optional[ResourceRequirements]: Average requirements or None if no history

"""

try:

if (

func_name not in self.history

or not self.history[func_name]

):

logger.debug(

"No history found for function {}".format(func_name)

```
)
```

```
return None
```

```
history = self.history[func_name]
```

```
avg_requirements = ResourceRequirements(
```

```
    cpu_memory_bytes=int(
```

```
        sum(r.cpu_memory_bytes for r in history)
```

```
        / len(history)
```

```
    ),
```

```
    gpu_memory_bytes=int(
```

```
        sum(r.gpu_memory_bytes for r in history)
```

```
        / len(history)
```

```
    ),
```

```
    requires_gpu=any(r.requires_gpu for r in history),
```

```
    estimated_runtime_seconds=sum(
```

```
        r.estimated_runtime_seconds for r in history
```

```
    )
```

```
    / len(history),
```

```
)
```

```
logger.debug(
```

```
    "Retrieved historical requirements for {}: {}".format(
```

```
        func_name,
```

```
        avg_requirements,
```

```
)
```

```
return avg_requirements
```

except Exception as e:

```
    logger.exception(
        "Error retrieving historical requirements for {}: {}".format(
            func_name,
            str(e),
        )
    )
    return None
```

def _estimate_object_size(self, obj: Any) -> int:

"""

Estimates the memory size of a Python object.

Args:

obj: Any Python object

Returns:

int: Estimated size in bytes

"""

try:

```
    with logger.contextualize(object_type=type(obj).__name__):
```

```
        if isinstance(obj, (np.ndarray, torch.Tensor)):
```

```
            size = obj.nbytes
```

```
            logger.debug("Array/Tensor size: {} bytes".format(size))
```

```
            return size
```

```
        elif isinstance(obj, (list, tuple, set, dict)):
```



```

size = sys.getsizeof(obj) + sum(
    self._estimate_object_size(item)
    for item in obj
)
logger.debug("Container size: {} bytes", size)
return size

```

```

size = sys.getsizeof(obj)
logger.debug("Object size: {} bytes", size)
return size

```

```

except Exception as e:
    logger.exception("Error estimating object size: {}", e)
    return 0

```

```

def _analyze_function_memory(
    self, func: Callable, *args, **kwargs

```

```

) -> ResourceRequirements:

```

```

"""

```

Analyzes function's memory requirements based on its arguments and source code.

Args:

func: Function to analyze

*args: Function arguments

**kwargs: Function keyword arguments

Returns:

ResourceRequirements: Predicted resource requirements

"""

```
logger.info(
    "Analyzing memory requirements for function: {}",
    func.__name__,
)
```

```
with logger.contextualize(function=func.__name__):
```

```
    try:
```

```
        # Estimate memory for input arguments
```

```
        args_size = sum(
            self._estimate_object_size(arg) for arg in args
        )
```

```
        kwargs_size = sum(
            self._estimate_object_size(value)
            for value in kwargs.values()
        )
```

```
        logger.debug("Arguments size: {} bytes", args_size)
```

```
        logger.debug(
            "Keyword arguments size: {} bytes", kwargs_size
        )
```

```
        # Analyze function source code for GPU operations
```

```
        source = inspect.getsource(func)
```

```
        requires_gpu = any(
```

```

        keyword in source

        for keyword in ["torch.cuda", "gpu", "GPU"]
    )

    logger.debug(

        "GPU requirement detected: {}", requires_gpu
    )


# Estimate GPU memory if needed (using a heuristic based on input size)
gpu_memory = args_size * 2 if requires_gpu else 0


# Add buffer for intermediate computations (50% of input size)
cpu_memory = (args_size + kwargs_size) * 1.5


requirements = ResourceRequirements(
    cpu_memory_bytes=int(cpu_memory),
    gpu_memory_bytes=int(gpu_memory),
    requires_gpu=requires_gpu,
    estimated_runtime_seconds=1.0,
)


logger.info(
    "Memory requirements analysis complete: {}",
    requirements,
)

return requirements

```

```
except Exception as e:
```

```
    logger.exception(  
        "Error analyzing function memory requirements: {}",  
        e,  
    )  
    raise
```

```
class ResourceManager:
```

```
    """Manages resource allocation for function execution."""
```

```
    def __init__(self):
```

```
        self.predictor = MemoryPredictor()
```

```
        self.reserved_cpu_memory = 0
```

```
        self.reserved_gpu_memory = {}
```

```
        logger.info("Initialized ResourceManager")
```

```
        # Log initial system resources
```

```
        self._log_system_resources()
```

```
    def _log_system_resources(self):
```

```
        """Logs current system resource status."""
```

```
        cpu_info = {
```

```
            "total_cpu_memory": psutil.virtual_memory().total,
```

```
            "available_cpu_memory": psutil.virtual_memory().available,
```

```
            "cpu_percent": psutil.cpu_percent(interval=1),
```

```
"cpu_count": psutil.cpu_count(),  
}
```

```
gpu_info = {}
```

```
if torch.cuda.is_available():
```

```
    for i in range(torch.cuda.device_count()):
```

```
        gpu_info[f"gpu_{i}"] = {
```

```
            "name": torch.cuda.get_device_name(i),
```

```
            "total_memory": torch.cuda.get_device_properties(  
                i
```

```
            ).total_memory,  
            "allocated_memory": torch.cuda.memory_allocated(  
                i
```

```
            ),  
        }  
}
```

```
logger.info("System resources - CPU: {}", cpu_info)
```

```
if gpu_info:
```

```
    logger.info("System resources - GPU: {}", gpu_info)
```

```
def _get_available_resources(self) -> Tuple[int, Dict[int, int]]:
```

```
    """
```

```
    Gets available CPU and GPU memory.
```

```
    Returns:
```

```
    Tuple[int, Dict[int, int]]: Available CPU memory and GPU memory per device
```

```
"""
```

```
with logger.contextualize(operation="resource_check"):

    cpu_available = psutil.virtual_memory().available

    gpu_available = {}

    if torch.cuda.is_available():

        for i in range(torch.cuda.device_count()):

            gpu_available[i] = (

                torch.cuda.get_device_properties(

                    i

                ).total_memory

                - torch.cuda.memory_allocated(i)

            )

        logger.debug(

            "Available CPU memory: {} bytes", cpu_available

        )

        logger.debug("Available GPU memory: {}", gpu_available)

    return cpu_available, gpu_available


@retry(

    stop=stop_after_attempt(3),

    wait=wait_fixed(1),

    retry=retry_if_exception_type(Exception),

)
```

```

def allocate_resources(
    self, requirements: ResourceRequirements
) -> ResourceAllocation:
    """
    Allocates required resources for function execution.

    Args:
        requirements: ResourceRequirements object with predicted requirements

    Returns:
        ResourceAllocation: Allocated resources

    Raises:
        InsufficientResourcesError: If required resources are not available
    """
    with logger.contextualize(operation="resource_allocation"):
        logger.info(
            "Attempting to allocate resources: {}", requirements
        )

        cpu_available, gpu_available = (
            self._get_available_resources()
        )

        # Check CPU memory availability
        if requirements.cpu_memory_bytes > cpu_available:

```

```
logger.error(  
    "Insufficient CPU memory. Required: {}, Available: {}",  
    requirements.cpu_memory_bytes,  
    cpu_available,  
)  
  
raise InsufficientResourcesError(  
    "Insufficient CPU memory"  
)
```

Select GPU if required

gpu_id = None

if requirements.requires_gpu:

for gid, available in gpu_available.items():

if available >= requirements.gpu_memory_bytes:

gpu_id = gid

logger.info("Selected GPU {}", gpu_id)

break

else:

logger.error(
 "No GPU with sufficient memory. Required: {}",
 requirements.gpu_memory_bytes,
)

raise InsufficientResourcesError(
 "Insufficient GPU memory"
)


```
# Allocate CPU cores based on memory requirements
```

```
total_cores = psutil.cpu_count()
```

```
cores_needed = max(
```

```
    1,
```

```
    min(
```

```
        total_cores,
```

```
        requirements.cpu_memory_bytes
```

```
        // (cpu_available // total_cores),
```

```
    ),
```

```
)
```

```
allocated_cores = list(range(cores_needed))
```

```
allocation = ResourceAllocation(
```

```
    cpu_memory_reserved=requirements.cpu_memory_bytes,
```

```
    gpu_memory_reserved=(
```

```
        requirements.gpu_memory_bytes
```

```
        if gpu_id is not None
```

```
        else 0
```

```
    ),
```

```
    gpu_id=gpu_id,
```

```
    cpu_cores=allocated_cores,
```

```
    allocation_time=time.time(),
```

```
)
```

```
logger.success(
```

```
    "Resource allocation successful: {}", allocation
```

)

return allocation

```
def release_resources(self, allocation: ResourceAllocation):
```

```
    """Releases allocated resources."""
```

```
    with logger.contextualize(operation="resource_release"):
```

```
        try:
```

```
            self.reserved_cpu_memory -= (
                allocation.cpu_memory_reserved
```

```
            )
```

```
            if allocation.gpu_id is not None:
```

```
                self.reserved_gpu_memory[
                    allocation.gpu_id
                ] -= allocation.gpu_memory_reserved
```

```
            allocation_duration = (
                time.time() - allocation.allocation_time
            )
```

```
            logger.info(
                "Resources released. Duration: {:.2f}s, CPU: {} bytes, GPU: {} bytes",
                allocation_duration,
                allocation.cpu_memory_reserved,
                allocation.gpu_memory_reserved,
            )
```

```
        except Exception as e:
```

```
            logger.exception("Error releasing resources: {}", e)
```

raise

```
def with_resource_management(func: Callable) -> Callable:
```

```
    """
```

Decorator that handles resource prediction and allocation for a function.

Args:

func: Function to be wrapped

Returns:

Callable: Wrapped function with resource management

```
    """
```

```
    resource_manager = ResourceManager()
```

```
    @wraps(func)
```

```
    def wrapper(*args, **kwargs):
```

```
        with logger.contextualize(
```

```
            function=func.__name__, execution_id=time.time_ns()
```

```
        ):
```

```
            logger.info("Starting execution with resource management")
```

```
            start_time = time.time()
```

```
            try:
```

```
                # Check historical data first
```

```
                historical_requirements = resource_manager.predictor.get_historical_requirements(
```

```

    func.__name__
)

if historical_requirements is not None:

    logger.info(

        "Using historical requirements: {}",

        historical_requirements,

    )

    requirements = historical_requirements

else:

    # Predict new requirements

    requirements = resource_manager.predictor._analyze_function_memory(

        func, *args, **kwargs

    )


# Allocate resources

allocation = resource_manager.allocate_resources(

    requirements

)


# Execute function with allocated resources

if allocation.gpu_id is not None:

    with torch.cuda.device(allocation.gpu_id):

        result = func(*args, **kwargs)

else:

    result = func(*args, **kwargs)

```

```

# Update execution history

execution_time = time.time() - start_time

requirements.estimated_runtime_seconds = (

    execution_time

)

resource_manager.predictor.update_history(

    func.__name__, requirements

)


logger.success(

    "Function execution completed successfully in {:.2f}s",

    execution_time,

)


return result


except Exception as e:

    logger.exception(

        "Error during function execution: {}", e

    )

    raise


finally:

    if "allocation" in locals():

        resource_manager.release_resources(allocation)

```

```
return wrapper
```

```
if __name__ == "__main__":
```

```
    @with_resource_management
```

```
    def process_large_matrix(size: int = 10000) -> np.ndarray:
```

```
        # Create a large matrix and perform memory-intensive operations
```

```
        data = np.random.random((size, size))
```

```
        result = np.matmul(data, data.T) # Matrix multiplication
```

```
        # Additional operations to increase memory usage
```

```
        for _ in range(3):
```

```
            result = np.exp(result) # Element-wise exponential
```

```
            result = np.sin(result) # Element-wise sine
```

```
            result = result @ result # Another matrix multiplication
```

```
        return result
```

```
    # Process a 10k x 10k matrix with intensive operations
```

```
    result = process_large_matrix()
```

```
    print(f"Result shape: {result.shape}")
```

```
    print(f"Peak memory usage: {result.nbytes / 1e9:.2f} GB")
```