

```
import asyncio
```

```
import time
```

```
from typing import Any, Callable, List, Tuple
```

```
import ray
```

```
import torch
```

```
from loguru import logger
```

```
from ray import serve
```

```
# Initialize Ray and Ray Serve
```

```
ray.init(ignore_reinit_error=True)
```

```
serve.start(detached=True)
```

```
# GPU Utility Function
```

```
def detect_gpus() -> List[str]:
```

```
    """
```

```
    Detects available GPU IDs on the system using torch.cuda.
```

```
    Returns:
```

```
        List[str]: A list of detected GPU IDs.
```

```
    """
```

```
    gpu_ids = [f"gpu-{i}" for i in range(torch.cuda.device_count())]
```

```
    logger.info(f"Detected GPUs: {gpu_ids}")
```

```
    return gpu_ids
```

```
@serve.deployment(route_prefix="/execute_job", num_replicas=2)
```

```
class GPUJobExecutor:
```

```
    def __init__(self):
```

```
        """Initialize GPUJobExecutor with a dynamic GPU pool and intelligent scheduler."""
```

```
        self.gpu_pool = detect_gpus()
```

```
        self.available_gpus = self.gpu_pool.copy()
```

```
        self.health_status = {
```

```
            gpu: True for gpu in self.gpu_pool
```

```
        } # GPU health tracking
```

```
        logger.info(f"Initialized GPU pool: {self.gpu_pool}")
```

```
        asyncio.create_task(self.monitor_gpu_health())
```

```
    async def monitor_gpu_health(self):
```

```
        """
```

```
        Monitors GPU health periodically and marks GPUs as unavailable if issues are detected.
```

```
        """
```

```
        while True:
```

```
            for gpu in self.gpu_pool:
```

```
                if self._check_gpu_health(gpu):
```

```
                    self.health_status[gpu] = True
```

```
                else:
```

```
                    logger.warning(f"GPU {gpu} marked unhealthy")
```

```
                    self.health_status[gpu] = False
```

```
            await asyncio.sleep(5) # Check every 5 seconds
```

```
def _check_gpu_health(self, gpu_id: str) -> bool:
```

```
    """
```

Checks the health of a GPU based on simple metrics.

This could be extended with more sophisticated monitoring in production.

Args:

gpu_id (str): GPU ID to check.

Returns:

bool: True if healthy, False if not.

```
    """
```

```
    gpu_index = int(gpu_id.split("-")[-1])
```

```
    memory_used = torch.cuda.memory_allocated(gpu_index)
```

```
    memory_limit = torch.cuda.get_device_properties(
```

```
        gpu_index
```

```
    ).total_memory
```

```
    # Basic health check: GPU is considered unhealthy if memory usage exceeds 90%
```

```
    return memory_used < 0.9 * memory_limit
```

```
async def execute(
```

```
    self,
```

```
    job_fn: Callable[[], Any],
```

```
    priority: int,
```

```
    *args,
```

```
    **kwargs,
```

```
) -> Any:
```

"""

Executes a job function with an available GPU in the pool.

Args:

job_fn (Callable): The job function to execute.

priority (int): Priority level for the job (1 is highest priority).

*args: Positional arguments for the job function.

**kwargs: Keyword arguments for the job function.

Returns:

Any: The result of the job function execution.

"""

```
if not self.available_gpus:
```

```
    logger.warning(
```

```
        "No GPUs available. Job will wait for an available GPU."
```

```
    )
```

```
    await self._wait_for_available_gpu()
```

```
gpu_id = self._allocate_gpu(priority)
```

```
logger.info(
```

```
    f"Allocating {gpu_id} for job execution with priority {priority}."
```

```
)
```

```
try:
```

```
    result = await job_fn(*args, **kwargs)
```

```
    logger.info(f"Job completed successfully on {gpu_id}.")
```

```
    return result
```

```
except Exception as e:
```

```
    logger.error(
```

```
        f"Error during job execution on {gpu_id}: {e}"
```

```
    )
```

```
    raise
```

```
finally:
```

```
    self._release_gpu(gpu_id)
```

```
async def _wait_for_available_gpu(self):
```

```
    """
```

```
    Waits until a GPU becomes available.
```

```
    """
```

```
    while not self.available_gpus:
```

```
        logger.info("Waiting for an available GPU...")
```

```
        await asyncio.sleep(1)
```

```
def _allocate_gpu(self, priority: int) -> str:
```

```
    """
```

```
    Allocates the first healthy GPU from the pool with priority handling.
```

```
    Returns:
```

```
        str: The allocated GPU ID.
```

```
    """
```

```
    healthy_gpus = [
```

```
        gpu
```

```

        for gpu in self.available_gpus
            if self.health_status[gpu]
    ]

    if healthy_gpus:

        gpu_id = healthy_gpus.pop(0)

        self.available_gpus.remove(gpu_id)

        logger.info(

            f"{gpu_id} allocated with priority {priority}."

        )

        return gpu_id

    else:

        logger.warning(

            "No healthy GPUs available; waiting for recovery."

        )

        time.sleep(1)

        return self._allocate_gpu(priority)

```

```

def _release_gpu(self, gpu_id: str):

```

```

    """

```

Releases a GPU back to the pool.

Args:

gpu_id (str): The GPU ID to release.

```

    """

```

```

    self.available_gpus.append(gpu_id)

```

```

    logger.info(f"{gpu_id} released back to pool.")

```

```

async def __call__(
    self,
    job_fn: Callable[[], Any],
    priority: int = 1,
    *args,
    **kwargs,
) -> Any:
    return await self.execute(job_fn, priority, *args, **kwargs)

```

@serve.deployment

class JobScheduler:

```

def __init__(self, executor_handle: Any):

```

```

    """

```

JobScheduler to manage job execution with fault tolerance, job retries, and scaling.

Args:

executor_handle (RayServeHandle): Handle to GPUJobExecutor deployment.

```

    """

```

```

self.executor_handle = executor_handle

```

```

self.max_retries = 3

```

```

self.job_queue: List[Tuple[int, Callable]] = (

```

```

    []

```

```

) # Prioritized job queue

```

```

self.min_executors = 2 # Minimum number of executors running

```

```
self.max_executors = 6 # Maximum number of executors allowed
```

```
asyncio.create_task(self.scale_cluster())
```

```
async def schedule_job(
```

```
    self,
```

```
    job_fn: Callable[[], Any],
```

```
    priority: int = 1,
```

```
    *args,
```

```
    **kwargs,
```

```
) -> Any:
```

```
    """
```

```
    Schedule and execute a job with retries in case of failure.
```

Args:

job_fn (Callable): The job function to execute.

priority (int): Priority level for the job (1 is highest priority).

*args: Positional arguments for the job function.

**kwargs: Keyword arguments for the job function.

Returns:

Any: The result of the job function execution.

```
    """
```

```
self.job_queue.append((priority, job_fn))
```

```
self.job_queue.sort(key=lambda x: x[0]) # Sort by priority
```

```
for attempt in range(self.max_retries):
```


try:

```
_, next_job = self.job_queue.pop(0)

result = await self.executor_handle.execute.remote(
    next_job, priority, *args, **kwargs
)

logger.info(
    f"Job completed successfully on attempt {attempt + 1}"
)

return result
```

except Exception as e:

```
logger.warning(
    f"Job failed on attempt {attempt + 1}: {e}"
)

if attempt == self.max_retries - 1:
    logger.error(
        "Max retries reached. Job failed permanently."
    )
    raise

await asyncio.sleep(2**attempt) # Exponential backoff

logger.info(
    f"Retrying job execution, attempt {attempt + 2}"
)
```

async def scale_cluster(self):

"""

Dynamically scales GPU nodes in the cluster based on load.

"""

while True:

 queue_size = len(self.job_queue)

 num_executors = len(

 await self.executor_handle.list_replicas.remote()

)

if queue_size > 5 and num_executors < self.max_executors:

 logger.info("Scaling up resources for high load.")

 new_executor = (

 GPUJobExecutor.bind()

) # Create a new executor

 ray.get(

 new_executor.deploy()

) # Deploy additional executor

elif (

 queue_size == 0 and num_executors > self.min_executors

):

 logger.info("Scaling down resources for low load.")

 await self.executor_handle.remove_replicas.remote(

 1

) # Remove an executor

 await asyncio.sleep(

 10

) # Adjust scaling check frequency as needed

```

# # Deployment

# gpu_executor = GPUJobExecutor.bind()

# job_scheduler = JobScheduler.bind(gpu_executor)


# # Example job function

# async def example_job_fn():

#     await asyncio.sleep(2) # Simulates job duration

#     return "Job completed"


# # Usage: Send request to job scheduler

# async def run():

#     job_handle = job_scheduler.schedule_job.remote(example_job_fn, priority=1)

#     result = await ray.get(job_handle)

#     print(result)


# # Deployments are initialized via Ray Serve:

# serve.run(run)


def gpu_scheduler(
    function: Callable[[[], Any], Any], priority: int = 1, *args, **kwargs
) -> Any:
    """

    Schedule and execute a function on available GPUs using Ray and Ray Serve.

```

Args:

function (Callable[[], Any]): The function to execute on GPU

priority (int, optional): Priority level for job scheduling. Lower numbers indicate higher priority.

Defaults to 1.

*args: Additional positional arguments to pass to the function

**kwargs: Additional keyword arguments to pass to the function

Returns:

Any: Result of the executed function

Raises:

RuntimeError: If there is an error initializing the GPU executor or job scheduler

TimeoutError: If the job execution times out

Exception: If there is an error executing the function

"""

try:

Initialize GPU executor and job scheduler

gpu_executor = GPUJobExecutor.bind()

job_scheduler = JobScheduler.bind(gpu_executor)

Schedule and execute the job

job_handle = job_scheduler.schedule_job.remote(

function, priority, *args, **kwargs

)

result = ray.get(job_handle, timeout=3600) # 1 hour timeout

```
logger.info(  
    f"Job completed successfully with result: {result}"  
)  
return result
```

```
except TimeoutError as e:
```

```
    logger.error(f"Job execution timed out: {e}")  
    raise TimeoutError(  
        "GPU job execution timed out after 1 hour"  
    ) from e
```

```
except Exception as e:
```

```
    logger.error(f"Error executing GPU job: {e}")  
    raise RuntimeError(  
        f"Failed to execute GPU job: {str(e)}"  
    ) from e
```