

```
import asyncio
```

```
import concurrent.futures
```

```
import logging
```

```
from typing import List, Tuple
```

```
import torch
```

```
from termcolor import colored
```

```
from transformers import (
```

```
    AutoModelForCausalLM,
```

```
    AutoTokenizer,
```

```
    BitsAndBytesConfig,
```

```
)
```

```
from swarm_models.base_llm import BaseLLM
```

```
class HuggingfaceLLM(BaseLLM):
```

```
    """
```

```
    A class for running inference on a given model.
```

```
    Attributes:
```

```
        model_id (str): The ID of the model.
```

```
        device (str): The device to run the model on (either 'cuda' or 'cpu').
```

```
        max_length (int): The maximum length of the output sequence.
```

```
        quantize (bool, optional): Whether to use quantization. Defaults to False.
```

```
        quantization_config (dict, optional): The configuration for quantization.
```

verbose (bool, optional): Whether to print verbose logs. Defaults to False.

logger (logging.Logger, optional): The logger to use. Defaults to a basic logger.

Methods:

run(task: str, max_length: int = 500) -> str:

Generate a response based on the prompt text.

__call__(task: str, max_length: int = 500) -> str:

Generate a response based on the prompt text.

save_model(path: str):

Save the model to a given path.

gpu_available() -> bool:

Check if GPU is available.

memory_consumption() -> dict:

Get the memory consumption of the GPU.

print_dashboard(task: str):

Print dashboard.

set_device(device: str):

Changes the device used for inference.

set_max_length(max_length: int):

Set max_length.

set_verbose(verbose: bool):

Set verbose.

set_distributed(distributed: bool):

Set distributed.

set_decoding(decoding: bool):

Set decoding.

set_max_workers(max_workers: int):

Set max_workers.

set_repitition_penalty(repitition_penalty: float):

Set repitition_penalty.

set_no_repeat_ngram_size(no_repeat_ngram_size: int):

Set no_repeat_ngram_size.

set_temperature(temperature: float):

Set temperature.

set_top_k(top_k: int):

Set top_k.

`set_top_p(top_p: float):`

Set top_p.

`set_quantize(quantize: bool):`

Set quantize.

`set_quantization_config(quantization_config: dict):`

Set quantization_config.

`set_model_id(model_id: str):`

Set model_id.

`set_model(model):`

Set model.

`set_tokenizer(tokenizer):`

Set tokenizer.

`set_logger(logger):`

Set logger.

Examples:

```
>>> llm = HuggingfaceLLM(  
...     model_id="EleutherAI/gpt-neo-2.7B",  
...     device="cuda",
```

```

...     max_length=500,
...     quantize=True,
...     quantization_config={
...         "load_in_4bit": True,
...         "bnb_4bit_use_double_quant": True,
...         "bnb_4bit_quant_type": "nf4",
...         "bnb_4bit_compute_dtype": torch.bfloat16,
...     },
... )

>>> llm("Generate a 10,000 word blog on mental clarity and the benefits of meditation.")

'Generate a 10,000 word
'''

```

```

def __init__(
    self,
    model_id: str,
    device: str = None,
    max_length: int = 500,
    quantize: bool = False,
    quantization_config: dict = None,
    verbose=False,
    distributed=False,
    decoding=False,
    max_workers: int = 5,
    repetition_penalty: float = 1.3,
    no_repeat_ngram_size: int = 5,

```

```
temperature: float = 0.7,

top_k: int = 40,

top_p: float = 0.8,

dtype=torch.bfloat16,

*args,

**kwargs,

):

    super().__init__(*args, **kwargs)

    self.logger = logging.getLogger(__name__)

    self.device = (

        device

        if device

        else ("cuda" if torch.cuda.is_available() else "cpu")

    )

    self.model_id = model_id

    self.max_length = max_length

    self.verbose = verbose

    self.distributed = distributed

    self.decoding = decoding

    self.quantize = quantize

    self.quantization_config = quantization_config

    self.max_workers = max_workers

    self.repitition_penalty = repitition_penalty

    self.no_repeat_ngram_size = no_repeat_ngram_size

    self.temperature = temperature

    self.top_k = top_k
```

```
self.top_p = top_p
```

```
self.dtype = dtype
```

```
if self.distributed:
```

```
    assert (
```

```
        torch.cuda.device_count() > 1
```

```
    ), "You need more than 1 gpu for distributed processing"
```

```
bnb_config = None
```

```
if quantize:
```

```
    if not quantization_config:
```

```
        quantization_config = {
```

```
            "load_in_4bit": True,
```

```
            "bnb_4bit_use_double_quant": True,
```

```
            "bnb_4bit_quant_type": "nf4",
```

```
            "bnb_4bit_compute_dtype": dtype,
```

```
        }
```

```
    bnb_config = BitsAndBytesConfig(**quantization_config)
```

```
self.tokenizer = AutoTokenizer.from_pretrained(self.model_id)
```

```
if quantize:
```

```
    self.model = AutoModelForCausalLM.from_pretrained(
```

```
        self.model_id,
```

```
        quantization_config=bnb_config,
```

```
        *args,
```

```
    **kwargs,
```

```
)
```

```
else:
```

```
    self.model = AutoModelForCausalLM.from_pretrained(
```

```
        self.model_id, *args, **kwargs
```

```
    ).to(self.device)
```

```
def print_error(self, error: str):
```

```
    """Print error"""
```

```
    print(colored(f"Error: {error}", "red"))
```

```
async def async_run(self, task: str):
```

```
    """Asynchronous generate text for a given prompt"""
```

```
    return await asyncio.to_thread(self.run, task)
```

```
def concurrent_run(self, tasks: List[str], max_workers: int = 5):
```

```
    """Concurrently generate text for a list of prompts."""
```

```
    with concurrent.futures.ThreadPoolExecutor(
```

```
        max_workers=max_workers
```

```
    ) as executor:
```

```
        results = list(executor.map(self.run, tasks))
```

```
    return results
```

```
def run_batch(
```

```
    self, tasks_images: List[Tuple[str, str]]
```

```
) -> List[str]:
```



```
"""Process a batch of tasks and images"""
```

with concurrent.futures.ThreadPoolExecutor() as executor:

```
futures = [  
    executor.submit(self.run, task, img)  
    for task, img in tasks_images  
]  
  
results = [future.result() for future in futures]  
  
return results
```

```
def run(self, task: str, *args, **kwargs):
```

```
    """
```

Generate a response based on the prompt text.

Args:

- task (str): Text to prompt the model.
- max_length (int): Maximum length of the response.

Returns:

- Generated text (str).

```
    """
```

```
    try:
```

```
        inputs = self.tokenizer.encode(task, return_tensors="pt")
```

```
        if self.decoding:
```

```
            with torch.no_grad():
```

```
                for _ in range(self.max_length):
```

```
output_sequence = []
```

```
outputs = self.model.generate(  
    inputs,  
    max_length=len(inputs) + 1,  
    do_sample=True,  
)
```

```
output_tokens = outputs[0][-1]
```

```
output_sequence.append(output_tokens.item())
```

```
# print token in real-time
```

```
print(  
    self.tokenizer.decode(  
        [output_tokens],  
        skip_special_tokens=True,  
    ),  
    end="",  
    flush=True,  
)
```

```
inputs = outputs
```

```
else:
```

```
    with torch.no_grad():
```

```
        outputs = self.model.generate(  
            inputs,  
            max_length=self.max_length,  
            do_sample=True,
```

```

        *args,
        **kwargs,
    )

    return self.tokenizer.decode(
        outputs[0], skip_special_tokens=True
    )
except Exception as e:
    print(
        colored(
            (
                "HuggingfaceLLM could not generate text"
                f" because of error: {e}, try optimizing your"
                " arguments"
            ),
            "red",
        )
    )
    raise

```

```

def __call__(self, task: str, *args, **kwargs):
    return self.run(task, *args, **kwargs)

```

```

async def __call_async__(self, task: str, *args, **kwargs) -> str:
    """Call the model asynchronously"""
    return await self.run_async(task, *args, **kwargs)

```

```

def save_model(self, path: str):
    """Save the model to a given path"""
    self.model.save_pretrained(path)
    self.tokenizer.save_pretrained(path)

def gpu_available(self) -> bool:
    """Check if GPU is available"""
    return torch.cuda.is_available()

def memory_consumption(self) -> dict:
    """Get the memory consumption of the GPU"""
    if self.gpu_available():
        torch.cuda.synchronize()
        allocated = torch.cuda.memory_allocated()
        reserved = torch.cuda.memory_reserved()
        return {"allocated": allocated, "reserved": reserved}
    else:
        return {"error": "GPU not available"}

def print_dashboard(self, task: str):
    """Print dashboard"""

    dashboard = print(
        colored(
            f"""

```

HuggingfaceLLM Dashboard

Model Name: {self.model_id}

Tokenizer: {self.tokenizer}

Model MaxLength: {self.max_length}

Model Device: {self.device}

Model Quantization: {self.quantize}

Model Quantization Config: {self.quantization_config}

Model Verbose: {self.verbose}

Model Distributed: {self.distributed}

Model Decoding: {self.decoding}

Metadata:

Task Memory Consumption: {self.memory_consumption()}

GPU Available: {self.gpu_available()}

Task Environment:

Task: {task}

""",

"red",

)

)

```
print(dashboard)
```

```
def set_device(self, device):
```

```
    """
```

```
    Changes the device used for inference.
```

```
    Parameters
```

```
    -----
```

```
        device : str
```

```
        The new device to use for inference.
```

```
    """
```

```
    self.device = device
```

```
    if self.model is not None:
```

```
        self.model.to(self.device)
```

```
def set_max_length(self, max_length):
```

```
    """Set max_length"""
```

```
    self.max_length = max_length
```

```
def clear_chat_history(self):
```

```
    """Clear chat history"""
```

```
    self.chat_history = []
```

```
def set_verbose(self, verbose):
```

```
    """Set verbose"""
```

```
    self.verbose = verbose
```

```
def set_distributed(self, distributed):
```

```
    """Set distributed"""
```

```
    self.distributed = distributed
```

```
def set_decoding(self, decoding):
```

```
    """Set decoding"""
```

```
    self.decoding = decoding
```

```
def set_max_workers(self, max_workers):
```

```
    """Set max_workers"""
```

```
    self.max_workers = max_workers
```

```
def set_repitition_penalty(self, repitition_penalty):
```

```
    """Set repitition_penalty"""
```

```
    self.repitition_penalty = repitition_penalty
```

```
def set_no_repeat_ngram_size(self, no_repeat_ngram_size):
```

```
    """Set no_repeat_ngram_size"""
```

```
    self.no_repeat_ngram_size = no_repeat_ngram_size
```

```
def set_temperature(self, temperature):
```

```
    """Set temperature"""
```

```
    self.temperature = temperature
```

```
def set_top_k(self, top_k):
```

```
"""Set top_k"""
```

```
self.top_k = top_k
```

```
def set_top_p(self, top_p):
```

```
    """Set top_p"""
```

```
    self.top_p = top_p
```

```
def set_quantize(self, quantize):
```

```
    """Set quantize"""
```

```
    self.quantize = quantize
```

```
def set_quantization_config(self, quantization_config):
```

```
    """Set quantization_config"""
```

```
    self.quantization_config = quantization_config
```

```
def set_model_id(self, model_id):
```

```
    """Set model_id"""
```

```
    self.model_id = model_id
```

```
def set_model(self, model):
```

```
    """Set model"""
```

```
    self.model = model
```

```
def set_tokenizer(self, tokenizer):
```

```
    """Set tokenizer"""
```

```
    self.tokenizer = tokenizer
```



```
def set_logger(self, logger):
```

```
    """Set logger"""
```

```
    self.logger = logger
```