

```
from typing import Callable, Optional
```

```
import torch
```

```
from termcolor import colored
```

```
from transformers import AutoProcessor, LdeficsForVisionText2Text
```

```
from swarm_models.base_multimodal_model import BaseMultiModalModel
```

```
def autodetect_device():
```

```
    """
```

```
    Autodetects the device to use for inference.
```

```
    Returns
```

```
    -----
```

```
    str
```

```
        The device to use for inference.
```

```
    """
```

```
    return "cuda" if torch.cuda.is_available() else "cpu"
```

```
class Ldefics(BaseMultiModalModel):
```

```
    """
```

```
    A class for multimodal inference using pre-trained models from the Hugging Face Hub.
```

Attributes

device : str

The device to use for inference.

model_name : str, optional

The name of the pre-trained model model_name (default is "HuggingFaceM4/idefics-9b-instruct").

processor : transformers.PreTrainedProcessor

The pre-trained processor.

max_length : int

The maximum length of the generated text.

chat_history : list

The chat history.

Methods

infer(prompts, batched_mode=True)

Generates text based on the provided prompts.

chat(user_input)

Engages in a continuous bidirectional conversation based on the user input.

set_model_name(model_name)

Changes the model model_name.

set_device(device)

Changes the device used for inference.

set_max_length(max_length)

Changes the maximum length of the generated text.

```
clear_chat_history()
```

Clears the chat history.

```
# Usage
```

```
...
```

```
from swarm_models import idefics
```

```
model = idefics()
```

```
user_input = "User: What is in this image?"
```

```
https://upload.wikimedia.org/wikipedia/commons/8/86/Id%C3%A9fix.JPG
```

```
response = model.chat(user_input)
```

```
print(response)
```

```
user_input = "User: And who is that?"
```

```
https://static.wikia.nocookie.net/asterix/images/2/25/R22b.gif/revision/latest?cb=20110815073052
```

```
response = model.chat(user_input)
```

```
print(response)
```

```
model.set_model_name("new_model_name")
```

```
model.set_device("cpu")
```

```
model.set_max_length(200)
```

```
model.clear_chat_history()
```

```
...
```

"""

```
def __init__(
    self,
    model_name: Optional[
        str
    ] = "HuggingFaceM4/idefics-9b-instruct",
    device: Callable = autodetect_device,
    torch_dtype=torch.bfloat16,
    max_length: int = 100,
    batched_mode: bool = True,
    *args,
    **kwargs,
):
    # Initialize the parent class
    super().__init__(*args, **kwargs)

    self.model_name = model_name

    self.device = device

    self.max_length = max_length

    self.batched_mode = batched_mode

    self.chat_history = []

    self.device = (
        device
        if device
        else ("cuda" if torch.cuda.is_available() else "cpu")
    )
```

```
)

self.model = LdeficsForVisionText2Text.from_pretrained(
    model_name, torch_dtype=torch_dtype, *args, **kwargs
).to(self.device)
```

```
self.processor = AutoProcessor.from_pretrained(
    model_name, *args, **kwargs
)
```

```
def run(
    self, task: str = None, img: str = None, *args, **kwargs
) -> str:
```

```
    """
```

Generates text based on the provided prompts.

Parameters

```
-----
```

task : str

the task to perform

batched_mode : bool, optional

Whether to process the prompts in batched mode. If True, all prompts are processed together. If False, only the first prompt is processed (default is True).

Returns

```
-----
```

list

A list of generated text strings.

```
"""
```

```
try:
```

```
    inputs = (  
        self.processor(  
            task,  
            add_end_of_utterance_token=False,  
            return_tensors="pt",  
            *args,  
            **kwargs,  
        ).to(self.device)  
        if self.batched_mode  
        else self.processor(task, return_tensors="pt").to(  
            self.device  
        )  
    )
```

```
exit_condition = self.processor.tokenizer(  
    "<end_of_utterance>", add_special_tokens=False  
).input_ids
```

```
bad_words_ids = self.processor.tokenizer(  
    ["<image>", "<fake_token_around_image>"],  
    add_special_tokens=False,  
).input_ids
```

```

generated_ids = self.model.generate(
    **inputs,
    eos_token_id=exit_condition,
    bad_words_ids=bad_words_ids,
    max_length=self.max_length,
)
generated_text = self.processor.batch_decode(
    generated_ids, skip_special_tokens=True
)
return generated_text

```

except Exception as error:

```

print(
    colored(
        (
            "Error in"
            f" {self.__class__.__name__} pipeline:"
            f" {error}"
        ),
        "red",
    )
)

```

```
def set_model_name(self, model_name):
```

```
    """
```

Changes the model model_name.

Parameters

`model_name : str`

The name of the new pre-trained model `model_name`.

"""

```
self.model = IdeficsForVisionText2Text.from_pretrained(
```

```
    model_name, torch_dtype=torch.bfloat16
```

```
).to(self.device)
```

```
self.processor = AutoProcessor.from_pretrained(model_name)
```