

```
resource "aws_instance" "k8s_master" {

  ami          = "ami-080e1f13689e07408" # Example AMI, replace with a Kubernetes supported one

  instance_type = "t3.xlarge"

  key_name      = aws_key_pair.ssh_key.key_name

  subnet_id     = aws_subnet.main.id # Ensure this is the corrected subnet ID

  iam_instance_profile = aws_iam_instance_profile.ec2_instance_profile.name

  vpc_security_group_ids = [aws_security_group.k8s_master_sg.id]

  tags = {

    Name = "KubernetesMaster"

  }

  root_block_device {

    volume_size = 50 # Specify the size of the root volume in GiB

    volume_type = "gp2"

  }

  user_data = base64encode(<<-EOSH

#!/bin/bash

# Update the system

sudo su

sudo apt-get update -y

sudo apt-get install -y iproute-tc

sudo setenforce 0

sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config

sudo apt-get -y install vim git curl wget

sudo swapoff -a

sudo sed -i '/swap/d' /etc/fstab
```

```
sudo mount -a
```

```
free -h
```

```
# Set up required modules
```

```
cat <<EOC | sudo tee /etc/modules-load.d/k8s.conf >/dev/null
```

```
overlay
```

```
br_netfilter
```

```
EOC
```

```
sudo modprobe overlay
```

```
sudo modprobe br_netfilter
```

```
# sysctl params required by setup, params persist across reboots
```

```
sudo cat <<EOK | sudo tee /etc/sysctl.d/k8s.conf > /dev/null
```

```
net.bridge.bridge-nf-call-iptables = 1
```

```
net.bridge.bridge-nf-call-ip6tables = 1
```

```
net.ipv4.ip_forward = 1
```

```
EOK
```

```
# Apply sysctl params without reboot
```

```
sudo sysctl --system
```

```
sudo apt-get install -y lvm2
```

```
# Install crictl (the version should match with the Kubernetes version you are using)
```

```
curl
```

-L

<https://github.com/kubernetes-sigs/cri-tools/releases/download/v1.29.0/crictl-v1.29.0-linux-amd64.tar>

.gz --output crictl.tar.gz

sudo tar zxvf crictl.tar.gz -C /usr/local/bin

rm -f crictl.tar.gz

sudo apt install -y containerd

# Configure the containerd cgroup driver to systemd

mkdir -p /etc/containerd

containerd config default | sudo tee /etc/containerd/config.toml

sudo sed -i 's/SystemdCgroup = false/SystemdCgroup = true/' /etc/containerd/config.toml

# Ensure the kubelet.service.d directory exists

sudo mkdir -p /etc/systemd/system/kubelet.service.d

sudo cat <<EOKUBE > /etc/systemd/system/kubelet.service.d/0-containerd-cgroup-driver.conf

[Service]

Environment="KUBELET\_EXTRA\_ARGS=--cgroup-driver=systemd"

EOKUBE

curl -LO "https://dl.k8s.io/release/\$(curl -L -s

https://dl.k8s.io/release/stable.txt)/bin/linux/amd64/kubectl"

curl -LO "https://dl.k8s.io/release/\$(curl -L -s

https://dl.k8s.io/release/stable.txt)/bin/linux/amd64/kubectl.sha256"

echo "\$(cat kubectl.sha256) kubectl" | sha256sum --check kubectl.sha256

sudo install -o root -g root -m 0755 kubectl /usr/local/bin/kubectl

sudo apt-get install -y apt-transport-https ca-certificates curl gpg

```
sudo mkdir -p -m 755 /etc/apt/keyrings

curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.29/deb/Release.key | sudo gpg --dearmor -o
/etc/apt/keyrings/kubernetes-apt-keyring.gpg

echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg]
https://pkgs.k8s.io/core:/stable:/v1.29/deb/ /' | sudo tee /etc/apt/sources.list.d/kubernetes.list

sudo apt-get update

sudo apt-get install -y kubelet kubeadm kubectl

sudo apt-mark hold kubelet kubeadm kubectl

# Restart containerd to apply the configuration changes

sudo systemctl restart containerd

# Reload systemd, enable and start kubelet

sudo systemctl daemon-reload

sudo systemctl enable --now kubelet containerd


sudo kubeadm config images pull --cri-socket unix:///run/containerd/containerd.sock

# Initialize the Kubernetes cluster

PUBLIC_IP=$(hostname -I | awk '{print $1}')

sudo kubeadm init --apiserver-advertise-address=$PUBLIC_IP --pod-network-cidr=10.244.0.0/16
--ignore-preflight-errors=all --v=10


# Set up kubeconfig for the root user

mkdir -p $HOME/.kube

sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config

sudo chown $(id -u):$(id -g) $HOME/.kube/config

export KUBECONFIG=$HOME/.kube/config
```

```
# Example readiness check loop for the kube-system namespace
```

```
for i in {1..10}; do
```

```
    kubectl get pods -n kube-system && break || sleep 15
```

```
done
```

```
# Apply network plugin if not already applied (idempotent operation)
```

```
echo "Applying Flannel CNI plugin..."
```

```
aws s3 cp s3://swarmskube/cogv1m_deployment.yml /tmp/cogv1m_deployment.yml
```

```
aws s3 cp s3://swarmskube/qwenv1_deployment.yml /tmp/qwenv1_deployment.yml
```

```
aws s3 cp s3://swarmskube/cogv1m_service.yml /tmp/cogv1m_service.yml
```

```
aws s3 cp s3://swarmskube/qwenv1_service.yml /tmp/qwenv1_service.yml
```

```
aws s3 cp s3://swarmskube/hpa.yml /tmp/hpa.yml
```

```
aws s3 cp s3://swarmskube/kubeflannel.yml.yml /tmp/kubeflannel.yml.yml
```

```
aws s3 cp s3://swarmskube/hpa.yml /tmp/router.yml
```

```
kubectl apply -f /tmp/kubeflannel.yml
```

```
kubectl apply -f /tmp/cogv1m_deployment.yml
```

```
kubectl apply -f /tmp/qwenv1_deployment.yml
```

```
kubectl apply -f /tmp/cogv1m_service.yml
```

```
kubectl apply -f /tmp/qwenv1_service.yml
```

```
kubectl apply -f /tmp/hpa.yml
```

```
kubectl apply -f /tmp/router.yml
```

```
curl https://baltocdn.com/helm/signing.asc | sudo apt-key add -
```

```
echo "deb https://baltocdn.com/helm/stable/debian/ all main" | sudo tee
```

```
/etc/apt/sources.list.d/helm-stable-debian.list
```

```
sudo apt-get update
```

```
sudo apt-get install -y helm
```

```
helm repo add prometheus-community https://prometheus-community.github.io/helm-charts
```

```
helm repo add grafana https://grafana.github.io/helm-charts
```

```
helm repo update
```

```
helm install prometheus prometheus-community/prometheus --namespace monitoring  
--create-namespace
```

```
helm install grafana grafana/grafana --namespace monitoring --create-namespace
```

```
sudo kubeadm token create --print-join-command | sudo tee /tmp/k8s-join-command.sh
```

```
sudo apt-get install -y awscli
```

```
aws s3 cp /tmp/k8s-join-command.sh s3://swarmskube/k8s-join-command.sh
```

```
aws s3 cp $HOME/.kube/config s3://swarmskube/kubeconfig
```

```
EOSH
```

```
)
```

```
}
```

```
resource "aws_launch_template" "k8s_worker" {
```

```
    name_prefix = "k8s-worker-"
```

```
    image_id      = "ami-080e1f13689e07408" # Example AMI, replace with a Kubernetes supported  
one
```

```
    instance_type = "p3.2xlarge"
```

```
    key_name      = aws_key_pair.ssh_key.key_name
```

```
    iam_instance_profile {
```

```
        name = aws_iam_instance_profile.ec2_instance_profile.name
```

```
}
```

```
vpc_security_group_ids = [aws_security_group.k8s_worker_sg.id]
```

```
block_device_mappings {
```

```
    device_name = "/dev/sda1"
```

```
    ebs {
```

```
        volume_size      = 50
```

```
        volume_type      = "gp3"
```

```
        delete_on_termination = true
```

```
    }
```

```
}
```

```
user_data = base64encode(<<-EOF
```

```
    #!/bin/bash
```

```
    # Download and execute the Kubernetes join command securely
```

```
sudo su
```

```
sudo apt-get update -y
```

```
sudo apt-get install -y iproute-tc
```

```
sudo setenforce 0
```

```
sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
```

```
sudo apt-get -y install vim git curl wget
```

```
sudo swapoff -a
```

```
sudo sed -i '/swap/d' /etc/fstab
```

```
sudo mount -a
```

```
free -h
```

```
# Set up required modules
```

```
cat <<EOC | sudo tee /etc/modules-load.d/k8s.conf >/dev/null
```

```
overlay
```

```
br_netfilter
```

```
EOC
```

```
sudo modprobe overlay
```

```
sudo modprobe br_netfilter
```

```
# sysctl params required by setup, params persist across reboots
```

```
sudo cat <<EOK | sudo tee /etc/sysctl.d/k8s.conf > /dev/null
```

```
net.bridge.bridge-nf-call-iptables = 1
```

```
net.bridge.bridge-nf-call-ip6tables = 1
```

```
net.ipv4.ip_forward = 1
```

```
EOK
```

```
# Apply sysctl params without reboot
```

```
sudo sysctl --system
```

```
sudo apt-get install -y lvm2
```

```
# Install crictl (the version should match with the Kubernetes version you are using)
```

```
curl
```

-L



<https://github.com/kubernetes-sigs/cri-tools/releases/download/v1.29.0/crictl-v1.29.0-linux-amd64.tar>

.gz --output crictl.tar.gz

sudo tar zxvf crictl.tar.gz -C /usr/local/bin

rm -f crictl.tar.gz

sudo apt install -y containerd

# Configure the containerd cgroup driver to systemd

mkdir -p /etc/containerd

containerd config default | sudo tee /etc/containerd/config.toml

sudo sed -i 's/SystemdCgroup = false/SystemdCgroup = true/' /etc/containerd/config.toml

# Ensure the kubelet.service.d directory exists

sudo mkdir -p /etc/systemd/system/kubelet.service.d

sudo cat <<EOKUBE > /etc/systemd/system/kubelet.service.d/0-containerd-cgroup-driver.conf

[Service]

Environment="KUBELET\_EXTRA\_ARGS=--cgroup-driver=systemd"

EOKUBE

curl -LO "https://dl.k8s.io/release/\$(curl -L -s

https://dl.k8s.io/release/stable.txt)/bin/linux/amd64/kubectl"

curl -LO "https://dl.k8s.io/release/\$(curl -L -s

https://dl.k8s.io/release/stable.txt)/bin/linux/amd64/kubectl.sha256"

echo "\$(cat kubectl.sha256) kubectl" | sha256sum --check kubectl.sha256

sudo install -o root -g root -m 0755 kubectl /usr/local/bin/kubectl

sudo apt-get install -y apt-transport-https ca-certificates curl gpg

```
sudo mkdir -p -m 755 /etc/apt/keyrings

curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.29/deb/Release.key | sudo gpg --dearmor -o
/etc/apt/keyrings/kubernetes-apt-keyring.gpg

echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg]
https://pkgs.k8s.io/core:/stable:/v1.29/deb/ /' | sudo tee /etc/apt/sources.list.d/kubernetes.list

sudo apt-get update

sudo apt-get install -y kubelet kubeadm kubectl

sudo apt-mark hold kubelet kubeadm kubectl

# Restart containerd to apply the configuration changes

sudo systemctl restart containerd

# Reload systemd, enable and start kubelet

sudo systemctl daemon-reload

sudo systemctl enable --now kubelet containerd

sudo apt-get install -y awscli
```

```
while true; do

    if aws s3 ls "s3://swarmskube/k8s-join-command.sh"; then

        aws s3 cp s3://swarmskube/k8s-join-command.sh /tmp/k8s-join-command.sh

        chmod +x /tmp/k8s-join-command.sh

        /tmp/k8s-join-command.sh

        break

    else

        echo "Waiting for the master node to initialize..."

        sleep 30

    fi

done
```

EOF

)

tags = {

resource\_type = "instance"

Name = "KubernetesWorker"

}

}