

```
import asyncio
```

```
import os
```

```
import threading
```

```
from concurrent.futures import ThreadPoolExecutor
```

```
from dataclasses import dataclass
```

```
from multiprocessing import cpu_count
```

```
from typing import Any, List
```

```
import psutil
```

```
from swarms.structs.agent import Agent
```

```
from swarms.structs.omni_agent_types import AgentType
```

```
from swarms.utils.wrapper_clusterop import (
```

```
    exec_callable_with_clusterops,
```

```
)
```

```
def run_single_agent(agent: AgentType, task: str) -> Any:
```

```
    """Run a single agent synchronously"""
```

```
    return agent.run(task)
```

```
async def run_agent_async(
```

```
    agent: AgentType, task: str, executor: ThreadPoolExecutor
```

```
) -> Any:
```

```
    """
```

Run an agent asynchronously using a thread executor.

Args:

agent: Agent instance to run

task: Task string to execute

executor: ThreadPoolExecutor instance for handling CPU-bound operations

Returns:

Agent execution result

"""

```
loop = asyncio.get_event_loop()
```

```
return await loop.run_in_executor(
```

```
    executor, run_single_agent, agent, task
```

```
)
```

```
async def run_agents_concurrently_async(
```

```
    agents: List[AgentType], task: str, executor: ThreadPoolExecutor
```

```
) -> List[Any]:
```

"""

Run multiple agents concurrently using asyncio and thread executor.

Args:

agents: List of Agent instances to run concurrently

task: Task string to execute

executor: ThreadPoolExecutor for CPU-bound operations

Returns:

List of outputs from each agent

```
"""
```

```
results = await asyncio.gather(
    *(run_agent_async(agent, task, executor) for agent in agents)
)
return results
```

```
def run_agents_concurrently(
```

```
    agents: List[AgentType],
```

```
    task: str,
```

```
    batch_size: int = None,
```

```
    max_workers: int = None,
```

```
) -> List[Any]:
```

```
"""
```

Optimized concurrent agent runner using both uvloop and ThreadPoolExecutor.

Args:

agents: List of Agent instances to run concurrently

task: Task string to execute

batch_size: Number of agents to run in parallel in each batch (defaults to CPU count)

max_workers: Maximum number of threads in the executor (defaults to CPU count * 2)

Returns:

List of outputs from each agent

```
"""
```

```
# Optimize defaults based on system resources
```

```
cpu_cores = cpu_count()
```

```
batch_size = batch_size or cpu_cores
```

```
max_workers = max_workers or cpu_cores * 2
```

```
results = []
```

```
# Get or create event loop
```

```
try:
```

```
    loop = asyncio.get_event_loop()
```

```
except RuntimeError:
```

```
    loop = asyncio.new_event_loop()
```

```
    asyncio.set_event_loop(loop)
```

```
# Create a shared thread pool executor with optimal worker count
```

```
with ThreadPoolExecutor(max_workers=max_workers) as executor:
```

```
    # Process agents in batches
```

```
    for i in range(0, len(agents), batch_size):
```

```
        batch = agents[i : i + batch_size]
```

```
        batch_results = loop.run_until_complete(
```

```
            run_agents_concurrently_async(batch, task, executor)
```

```
        )
```

```
        results.extend(batch_results)
```

return results

```
def run_agents_concurrently_multiprocess(
```

```
    agents: List[Agent], task: str, batch_size: int = cpu_count()
```

```
) -> List[Any]:
```

```
    """
```

Manage and run multiple agents concurrently in batches, with optimized performance.

Args:

agents (List[Agent]): List of Agent instances to run concurrently.

task (str): The task string to execute by all agents.

batch_size (int, optional): Number of agents to run in parallel in each batch.

Defaults to the number of CPU cores.

Returns:

List[Any]: A list of outputs from each agent.

```
    """
```

```
    results = []
```

```
    loop = asyncio.get_event_loop()
```

```
    # Process agents in batches to avoid overwhelming system resources
```

```
    for i in range(0, len(agents), batch_size):
```

```
        batch = agents[i : i + batch_size]
```

```
        batch_results = loop.run_until_complete(
```

```
            run_agents_concurrently_async(batch, task)
```

)

results.extend(batch_results)

return results

def run_agents_sequentially(
 agents: List[AgentType], task: str

) -> List[Any]:

"""

Run multiple agents sequentially for baseline comparison.

Args:

agents: List of Agent instances to run

task: Task string to execute

Returns:

List of outputs from each agent

"""

return [run_single_agent(agent, task) for agent in agents]

def run_agents_with_different_tasks(
 agent_task_pairs: List[tuple[AgentType, str]],

batch_size: int = None,

max_workers: int = None,

) -> List[Any]:

"""

Run multiple agents with different tasks concurrently.

Args:

agent_task_pairs: List of (agent, task) tuples

batch_size: Number of agents to run in parallel

max_workers: Maximum number of threads

Returns:

List of outputs from each agent

"""

async def run_pair_async(

pair: tuple[AgentType, str], executor: ThreadPoolExecutor

) -> Any:

agent, task = pair

return await run_agent_async(agent, task, executor)

cpu_cores = cpu_count()

batch_size = batch_size or cpu_cores

max_workers = max_workers or cpu_cores * 2

results = []

try:

loop = asyncio.get_event_loop()

except RuntimeError:

loop = asyncio.new_event_loop()

asyncio.set_event_loop(loop)

with ThreadPoolExecutor(max_workers=max_workers) as executor:

for i in range(0, len(agent_task_pairs), batch_size):

batch = agent_task_pairs[i : i + batch_size]

batch_results = loop.run_until_complete(

asyncio.gather(

*(

run_pair_async(pair, executor)

for pair in batch

)

)

)

results.extend(batch_results)

return results

async def run_agent_with_timeout(

agent: AgentType,

task: str,

timeout: float,

executor: ThreadPoolExecutor,

) -> Any:

"""

Run an agent with a timeout limit.

Args:

agent: Agent instance to run

task: Task string to execute

timeout: Timeout in seconds

executor: ThreadPoolExecutor instance

Returns:

Agent execution result or None if timeout occurs

"""

try:

return await asyncio.wait_for(

run_agent_async(agent, task, executor), timeout=timeout

)

except asyncio.TimeoutError:

return None

def run_agents_with_timeout(

agents: List[AgentType],

task: str,

timeout: float,

batch_size: int = None,

max_workers: int = None,

) -> List[Any]:

"""

Run multiple agents concurrently with a timeout for each agent.

Args:

agents: List of Agent instances

task: Task string to execute

timeout: Timeout in seconds for each agent

batch_size: Number of agents to run in parallel

max_workers: Maximum number of threads

Returns:

List of outputs (None for timed out agents)

"""

cpu_cores = cpu_count()

batch_size = batch_size or cpu_cores

max_workers = max_workers or cpu_cores * 2

results = []

try:

loop = asyncio.get_event_loop()

except RuntimeError:

loop = asyncio.new_event_loop()

asyncio.set_event_loop(loop)

with ThreadPoolExecutor(max_workers=max_workers) as executor:

```

for i in range(0, len(agents), batch_size):

    batch = agents[i : i + batch_size]

    batch_results = loop.run_until_complete(

        asyncio.gather(

            *(

                run_agent_with_timeout(

                    agent, task, timeout, executor

                )

                for agent in batch

            )

        )

    )

    results.extend(batch_results)

return results

```

```
@dataclass
```

```
class ResourceMetrics:
```

```
    cpu_percent: float
```

```
    memory_percent: float
```

```
    active_threads: int
```

```
def get_system_metrics() -> ResourceMetrics:
```

```
    """Get current system resource usage"""

```

```

return ResourceMetrics(
    cpu_percent=psutil.cpu_percent(),
    memory_percent=psutil.virtual_memory().percent,
    active_threads=threading.active_count(),
)

```

```

def run_agents_with_resource_monitoring(

```

```

    agents: List[AgentType],
    task: str,
    cpu_threshold: float = 90.0,
    memory_threshold: float = 90.0,
    check_interval: float = 1.0,

```

```

) -> List[Any]:

```

```

    """

```

Run agents with system resource monitoring and adaptive batch sizing.

Args:

agents: List of Agent instances

task: Task string to execute

cpu_threshold: Max CPU usage percentage

memory_threshold: Max memory usage percentage

check_interval: Resource check interval in seconds

Returns:

List of outputs from each agent

```
"""
```

```
async def monitor_resources():
```

```
    while True:
```

```
        metrics = get_system_metrics()
```

```
        if (
```

```
            metrics.cpu_percent > cpu_threshold
```

```
            or metrics.memory_percent > memory_threshold
```

```
        ):
```

```
            # Reduce batch size or pause execution
```

```
            pass
```

```
        await asyncio.sleep(check_interval)
```

```
# Implementation details...
```

```
def _run_agents_with_tasks_concurrently(
```

```
    agents: List[AgentType],
```

```
    tasks: List[str] = [],
```

```
    batch_size: int = None,
```

```
    max_workers: int = None,
```

```
) -> List[Any]:
```

```
"""
```

```
Run multiple agents with corresponding tasks concurrently.
```

```
Args:
```

agents: List of Agent instances to run

tasks: List of task strings to execute

batch_size: Number of agents to run in parallel

max_workers: Maximum number of threads

Returns:

List of outputs from each agent

```
"""
```

```
if len(agents) != len(tasks):
```

```
    raise ValueError(
```

```
        "The number of agents must match the number of tasks."
```

```
)
```

```
cpu_cores = os.cpu_count()
```

```
batch_size = batch_size or cpu_cores
```

```
max_workers = max_workers or cpu_cores * 2
```

```
results = []
```

```
try:
```

```
    loop = asyncio.get_event_loop()
```

```
except RuntimeError:
```

```
    loop = asyncio.new_event_loop()
```

```
    asyncio.set_event_loop(loop)
```

```
async def run_agent_task_pair(
```

```
    agent: AgentType, task: str, executor: ThreadPoolExecutor
```

) -> Any:

```
    return await run_agent_async(agent, task, executor)
```

```
with ThreadPoolExecutor(max_workers=max_workers) as executor:
```

```
    for i in range(0, len(agents), batch_size):
```

```
        batch_agents = agents[i : i + batch_size]
```

```
        batch_tasks = tasks[i : i + batch_size]
```

```
        batch_results = loop.run_until_complete(
```

```
            asyncio.gather(
```

```
                *(
```

```
                    run_agent_task_pair(agent, task, executor)
```

```
                    for agent, task in zip(
```

```
                        batch_agents, batch_tasks
```

```
                    )
```

```
                )
```

```
            )
```

```
        )
```

```
        results.extend(batch_results)
```

```
    return results
```

```
def run_agents_with_tasks_concurrently(
```

```
    agents: List[AgentType],
```

```
    tasks: List[str] = [],
```

```
    batch_size: int = None,
```

```
max_workers: int = None,  
device: str = "cpu",  
device_id: int = 1,  
all_cores: bool = True,  
no_clusterops: bool = False,  
) -> List[Any]:
```

```
"""
```

Executes a list of agents with their corresponding tasks concurrently on a specified device.

This function orchestrates the concurrent execution of a list of agents with their respective tasks on a specified device, either CPU or GPU. It leverages the ``exec_callable_with_clusterops`` function to manage the execution on the specified device.

Args:

`agents (List[AgentType])`: A list of Agent instances or callable functions to execute concurrently.

`tasks (List[str], optional)`: A list of task strings to execute for each agent. Defaults to an empty list.

`batch_size (int, optional)`: The number of agents to run in parallel. Defaults to None.

`max_workers (int, optional)`: The maximum number of threads to use for execution. Defaults to None.

`device (str, optional)`: The device to use for execution. Defaults to "cpu".

`device_id (int, optional)`: The ID of the GPU to use if device is set to "gpu". Defaults to 0.

`all_cores (bool, optional)`: If True, uses all available CPU cores. Defaults to True.

Returns:

`List[Any]`: A list of outputs from each agent execution.


```
"""
```

```
# Make the first agent not use the ifrs
```

```
if no_clusterops:
```

```
    return _run_agents_with_tasks_concurrently(
        agents, tasks, batch_size, max_workers
    )
```

```
else:
```

```
    return exec_callable_with_clusterops(
        device,
        device_id,
        all_cores,
        _run_agents_with_tasks_concurrently,
        agents,
        tasks,
        batch_size,
        max_workers,
    )
```

```
# # Example usage:
```

```
# # Initialize your agents with the same model to avoid re-creating it
```

```
# agents = [
```

```
#     Agent(
```

```
#         agent_name=f"Financial-Analysis-Agent_parallel_swarm{i}",
```

```
#         system_prompt=FINANCIAL_AGENT_SYS_PROMPT,
```

```

#     llm=model,
#     max_loops=1,
#     autosave=True,
#     dashboard=False,
#     verbose=False,
#     dynamic_temperature_enabled=False,
#     saved_state_path=f"finance_agent_{i}.json",
#     user_name="swarms_corp",
#     retry_attempts=1,
#     context_length=200000,
#     return_step_meta=False,
# )
# for i in range(5) # Assuming you want 10 agents
# ]

# task = "How can I establish a ROTH IRA to buy stocks and get a tax break? What are the criteria"
# outputs = run_agents_concurrently(agents, task)

# for i, output in enumerate(outputs):
#     print(f"Output from agent {i+1}:\n{output}")

```