

```
"""Sampling parameters for text generation."""
```

```
from enum import IntEnum
```

```
from functools import cached_property
```

```
from typing import Callable, List, Optional, Union
```

```
import torch
```

```
_SAMPLING_EPS = 1e-5
```

```
class SamplingType(IntEnum):
```

```
    GREEDY = 0
```

```
    RANDOM = 1
```

```
    BEAM = 2
```

```
LogitsProcessor = Callable[[List[int], torch.Tensor], torch.Tensor]
```

```
"""LogitsProcessor is a function that takes a list of previously generated  
tokens and a tensor of the logits for the next token, and returns a modified  
tensor of logits to sample from."""
```

```
class SamplingParams:
```

```
    """Sampling parameters for text generation.
```

Overall, we follow the sampling parameters from the OpenAI text completion API (<https://platform.openai.com/docs/api-reference/completions/create>). In addition, we support beam search, which is not supported by OpenAI.

Args:

n: Number of output sequences to return for the given prompt.

best_of: Number of output sequences that are generated from the prompt.

From these `best_of` sequences, the top `n` sequences are returned.

`best_of` must be greater than or equal to `n`. This is treated as the beam width when `use_beam_search` is True. By default, `best_of` is set to `n`.

presence_penalty: Float that penalizes new tokens based on whether they appear in the generated text so far. Values > 0 encourage the model to use new tokens, while values < 0 encourage the model to repeat tokens.

frequency_penalty: Float that penalizes new tokens based on their frequency in the generated text so far. Values > 0 encourage the model to use new tokens, while values < 0 encourage the model to repeat tokens.

repetition_penalty: Float that penalizes new tokens based on whether they appear in the prompt and the generated text so far. Values > 1 encourage the model to use new tokens, while values < 1 encourage the model to repeat tokens.

temperature: Float that controls the randomness of the sampling. Lower values make the model more deterministic, while higher values make the model more random. Zero means greedy sampling.

`top_p`: Float that controls the cumulative probability of the top tokens to consider. Must be in (0, 1]. Set to 1 to consider all tokens.

`top_k`: Integer that controls the number of top tokens to consider. Set to -1 to consider all tokens.

`min_p`: Float that represents the minimum probability for a token to be considered, relative to the probability of the most likely token. Must be in [0, 1]. Set to 0 to disable this.

`use_beam_search`: Whether to use beam search instead of sampling.

`length_penalty`: Float that penalizes sequences based on their length. Used in beam search.

`early_stopping`: Controls the stopping condition for beam search. It accepts the following values: ``True``, where the generation stops as soon as there are ``best_of`` complete candidates; ``False``, where an heuristic is applied and the generation stops when is it very unlikely to find better candidates; ``"never"``, where the beam search procedure only stops when there cannot be better candidates (canonical beam search algorithm).

`stop`: List of strings that stop the generation when they are generated. The returned output will not contain the stop strings.

`stop_token_ids`: List of tokens that stop the generation when they are generated. The returned output will contain the stop tokens unless the stop tokens are special tokens.

`include_stop_str_in_output`: Whether to include the stop strings in output text. Defaults to False.

`ignore_eos`: Whether to ignore the EOS token and continue generating tokens after the EOS token is generated.

max_tokens: Maximum number of tokens to generate per output sequence.

logprobs: Number of log probabilities to return per output token.

Note that the implementation follows the OpenAI API: The return

result includes the log probabilities on the `logprobs` most likely

tokens, as well the chosen tokens. The API will always return the

log probability of the sampled token, so there may be up to

`logprobs+1` elements in the response.

prompt_logprobs: Number of log probabilities to return per prompt token.

skip_special_tokens: Whether to skip special tokens in the output.

spaces_between_special_tokens: Whether to add spaces between special

tokens in the output. Defaults to True.

logits_processors: List of functions that modify logits based on

previously generated tokens.

"""

```
def __init__(
    self,
    n: int = 1,
    best_of: Optional[int] = None,
    presence_penalty: float = 0.0,
    frequency_penalty: float = 0.0,
    repetition_penalty: float = 1.0,
    temperature: float = 1.0,
    top_p: float = 1.0,
    top_k: int = -1,
    min_p: float = 0.0,
```

use_beam_search: bool = False,
length_penalty: float = 1.0,
early_stopping: Union[bool, str] = False,
stop: Union[str, List[str], None] = None,
stop_token_ids: Optional[List[int]] = None,
include_stop_str_in_output: bool = False,
ignore_eos: bool = False,
max_tokens: Optional[int] = 16,
logprobs: Optional[int] = None,
prompt_logprobs: Optional[int] = None,
skip_special_tokens: bool = True,
spaces_between_special_tokens: bool = True,
logits_processors: Optional[List[LogitsProcessor]] = None,

) -> None:

self.n = n
self.best_of = best_of if best_of is not None else n
self.presence_penalty = presence_penalty
self.frequency_penalty = frequency_penalty
self.repetition_penalty = repetition_penalty
self.temperature = temperature
self.top_p = top_p
self.top_k = top_k
self.min_p = min_p
self.use_beam_search = use_beam_search
self.length_penalty = length_penalty
self.early_stopping = early_stopping

```
if stop is None:
    self.stop = []
elif isinstance(stop, str):
    self.stop = [stop]
else:
    self.stop = list(stop)
if stop_token_ids is None:
    self.stop_token_ids = []
else:
    self.stop_token_ids = list(stop_token_ids)
self.ignore_eos = ignore_eos
self.max_tokens = max_tokens
self.logprobs = logprobs
self.prompt_logprobs = prompt_logprobs
self.skip_special_tokens = skip_special_tokens
self.spaces_between_special_tokens = (
    spaces_between_special_tokens
)
self.logits_processors = logits_processors
self.include_stop_str_in_output = include_stop_str_in_output
self._verify_args()
if self.use_beam_search:
    self._verify_beam_search()
else:
    self._verify_non_beam_search()
    if self.temperature < _SAMPLING_EPS:
```

```
# Zero temperature means greedy sampling.
```

```
self.top_p = 1.0
```

```
self.top_k = -1
```

```
self.min_p = 0.0
```

```
self._verify_greedy_sampling()
```

```
def _verify_args(self) -> None:
```

```
    if self.n < 1:
```

```
        raise ValueError(f"n must be at least 1, got {self.n}.")
```

```
    if self.best_of < self.n:
```

```
        raise ValueError(
```

```
            "best_of must be greater than or equal to n, "
```

```
            f"got n={self.n} and best_of={self.best_of}."
```

```
        )
```

```
    if not -2.0 <= self.presence_penalty <= 2.0:
```

```
        raise ValueError(
```

```
            "presence_penalty must be in [-2, 2], got "
```

```
            f"{self.presence_penalty}."
```

```
        )
```

```
    if not -2.0 <= self.frequency_penalty <= 2.0:
```

```
        raise ValueError(
```

```
            "frequency_penalty must be in [-2, 2], got "
```

```
            f"{self.frequency_penalty}."
```

```
        )
```

```
    if not 0.0 < self.repetition_penalty <= 2.0:
```

```
        raise ValueError(
```

```

        "repetition_penalty must be in (0, 2], got "
        f"{self.repetition_penalty}."
    )

if self.temperature < 0.0:
    raise ValueError(
        "temperature must be non-negative, got"
        f" {self.temperature}."
    )

if not 0.0 < self.top_p <= 1.0:
    raise ValueError(
        f"top_p must be in (0, 1], got {self.top_p}."
    )

if self.top_k < -1 or self.top_k == 0:
    raise ValueError(
        "top_k must be -1 (disable), or at least 1, "
        f"got {self.top_k}."
    )

if not 0.0 <= self.min_p <= 1.0:
    raise ValueError(
        f"min_p must be in [0, 1], got {self.min_p}."
    )

if self.max_tokens is not None and self.max_tokens < 1:
    raise ValueError(
        "max_tokens must be at least 1, got"
        f" {self.max_tokens}."
    )

```



```

if self.logprobs is not None and self.logprobs < 0:

    raise ValueError(

        f"logprobs must be non-negative, got {self.logprobs}."

    )

if (

    self.prompt_logprobs is not None

    and self.prompt_logprobs < 0

):

    raise ValueError(

        "prompt_logprobs must be non-negative, got "

        f"{self.prompt_logprobs}."

    )

```

```

def _verify_beam_search(self) -> None:

```

```

    if self.best_of == 1:

        raise ValueError(

            "best_of must be greater than 1 when using beam "

            f"search. Got {self.best_of}."

        )

```

```

    if self.temperature > _SAMPLING_EPS:

```

```

        raise ValueError(

            "temperature must be 0 when using beam search."

        )

```

```

    if self.top_p < 1.0 - _SAMPLING_EPS:

```

```

        raise ValueError(

            "top_p must be 1 when using beam search."

        )

```

)

if self.top_k != -1:

raise ValueError(

"top_k must be -1 when using beam search."

)

if self.early_stopping not in [True, False, "never"]:

raise ValueError(

"early_stopping must be True, False, or 'never', "

f"got {self.early_stopping}."

)

def _verify_non_beam_search(self) -> None:

if self.early_stopping is not False:

raise ValueError(

"early_stopping is not effective and must be "

"False when not using beam search."

)

if (

self.length_penalty < 1.0 - _SAMPLING_EPS

or self.length_penalty > 1.0 + _SAMPLING_EPS

):

raise ValueError(

"length_penalty is not effective and must be the "

"default value of 1.0 when not using beam search."

)

```
def _verify_greedy_sampling(self) -> None:

    if self.best_of > 1:

        raise ValueError(

            "best_of must be 1 when using greedy sampling."

            f"Got {self.best_of}."

        )
```

```
@cached_property
```

```
def sampling_type(self) -> SamplingType:

    if self.use_beam_search:

        return SamplingType.BEAM

    if self.temperature < _SAMPLING_EPS:

        return SamplingType.GREEDY

    return SamplingType.RANDOM
```

```
def __repr__(self) -> str:

    return (

        f"SamplingParams(n={self.n}, "

        f"best_of={self.best_of}, "

        f"presence_penalty={self.presence_penalty}, "

        f"frequency_penalty={self.frequency_penalty}, "

        f"repetition_penalty={self.repetition_penalty}, "

        f"temperature={self.temperature}, "

        f"top_p={self.top_p}, "

        f"top_k={self.top_k}, "

        f"min_p={self.min_p}, "
```

```
f"use_beam_search={self.use_beam_search}, "  
f"length_penalty={self.length_penalty}, "  
f"early_stopping={self.early_stopping}, "  
f"stop={self.stop}, "  
f"stop_token_ids={self.stop_token_ids}, "  
f"include_stop_str_in_output={self.include_stop_str_in_output}, "  
f"ignore_eos={self.ignore_eos}, "  
f"max_tokens={self.max_tokens}, "  
f"logprobs={self.logprobs}, "  
f"prompt_logprobs={self.prompt_logprobs}, "  
f"skip_special_tokens={self.skip_special_tokens}, "  
"spaces_between_special_tokens="  
f"{self.spaces_between_special_tokens})"
```