

```
import logging

from unittest.mock import patch


import pytest

import torch


from swarm_models.huggingface import HuggingfaceLLM


# Fixture for the class instance

@pytest.fixture

def llm_instance():

    model_id = "NousResearch/Nous-Hermes-2-Vision-Alpha"

    instance = HuggingfaceLLM(model_id=model_id)

    return instance


# Test for instantiation and attributes

def test_llm_initialization(llm_instance):

    assert (

        llm_instance.model_id

        == "NousResearch/Nous-Hermes-2-Vision-Alpha"

    )

    assert llm_instance.max_length == 500

    # ... add more assertions for all default attributes
```

```

# Parameterized test for setting devices

@pytest.mark.parametrize("device", ["cpu", "cuda"])

def test_llm_set_device(llm_instance, device):

    llm_instance.set_device(device)

    assert llm_instance.device == device


# Test exception during initialization with a bad model_id

def test_llm_bad_model_initialization():

    with pytest.raises(Exception):

        HuggingfaceLLM(model_id="unknown-model")


# # Mocking the tokenizer and model to test run method

# @patch("swarms.models.huggingface.AutoTokenizer.from_pretrained")

# @patch(

#     "swarms.models.huggingface.AutoModelForCausalLM.from_pretrained"

# )

# def test_llm_run(mock_model, mock_tokenizer, llm_instance):

#     mock_model.return_value.generate.return_value = "mocked output"

#     mock_tokenizer.return_value.encode.return_value = "mocked input"

#     result = llm_instance.run("test task")

#     assert result == "mocked output"

```

```
# Async test (requires pytest-asyncio plugin)
```

```
@pytest.mark.asyncio
```

```
async def test_llm_run_async(llm_instance):
```

```
    result = await llm_instance.run_async("test task")
```

```
    assert isinstance(result, str)
```

```
# Test for checking GPU availability
```

```
def test_llm_gpu_availability(llm_instance):
```

```
    # Assuming the test is running on a machine where the GPU availability is known
```

```
    expected_result = torch.cuda.is_available()
```

```
    assert llm_instance.gpu_available() == expected_result
```

```
# Test for memory consumption reporting
```

```
def test_llm_memory_consumption(llm_instance):
```

```
    # Mocking torch.cuda functions for consistent results
```

```
    with patch("torch.cuda.memory_allocated", return_value=1024):
```

```
        with patch("torch.cuda.memory_reserved", return_value=2048):
```

```
            memory = llm_instance.memory_consumption()
```

```
            assert memory == {"allocated": 1024, "reserved": 2048}
```

```
# Test different initialization parameters
```

```
@pytest.mark.parametrize(
```

```
    "model_id, max_length",
```

```
[
    ("NousResearch/Nous-Hermes-2-Vision-Alpha", 100),
    ("microsoft/Orca-2-13b", 200),
    (
        "berkeley-nest/Starling-LM-7B-alpha",
        None,
    ), # None to check default behavior
],
)
```

```
def test_llm_initialization_params(model_id, max_length):
```

```
    if max_length:
```

```
        instance = HuggingfaceLLM(
            model_id=model_id, max_length=max_length
        )
```

```
        assert instance.max_length == max_length
```

```
    else:
```

```
        instance = HuggingfaceLLM(model_id=model_id)
```

```
        assert (
            instance.max_length == 500
```

```
        ) # Assuming 500 is the default max_length
```

```
# Test for setting an invalid device
```

```
def test_llm_set_invalid_device(llm_instance):
```

```
    with pytest.raises(ValueError):
```

```
        llm_instance.set_device("quantum_processor")
```

```
# Mocking external API call to test run method without network
```

```
@patch("swarms.models.huggingface.HuggingfaceLLM.run")
```

```
def test_llm_run_without_network(mock_run, llm_instance):
```

```
    mock_run.return_value = "mocked output"
```

```
    result = llm_instance.run("test task without network")
```

```
    assert result == "mocked output"
```

```
# Test handling of empty input for the run method
```

```
def test_llm_run_empty_input(llm_instance):
```

```
    with pytest.raises(ValueError):
```

```
        llm_instance.run("")
```

```
# Test the generation with a provided seed for reproducibility
```

```
@patch("swarms.models.huggingface.HuggingfaceLLM.run")
```

```
def test_llm_run_with_seed(mock_run, llm_instance):
```

```
    seed = 42
```

```
    llm_instance.set_seed(seed)
```

```
    # Assuming set_seed method affects the randomness in the model
```

```
    # You would typically ensure that setting the seed gives reproducible results
```

```
    mock_run.return_value = "mocked deterministic output"
```

```
    result = llm_instance.run("test task", seed=seed)
```

```
    assert result == "mocked deterministic output"
```

```
# Test the output length is as expected
```

```
@patch("swarms.models.huggingface.HuggingfaceLLM.run")
```

```
def test_llm_run_output_length(mock_run, llm_instance):
```

```
    input_text = "test task"
```

```
    llm_instance.max_length = 50 # set a max_length for the output
```

```
    mock_run.return_value = "mocked output" * 10 # some long text
```

```
    result = llm_instance.run(input_text)
```

```
    assert len(result.split()) <= llm_instance.max_length
```

```
# Test the tokenizer handling special tokens correctly
```

```
@patch("swarms.models.huggingface.HuggingfaceLLM._tokenizer.encode")
```

```
@patch("swarms.models.huggingface.HuggingfaceLLM._tokenizer.decode")
```

```
def test_llm_tokenizer_special_tokens(
```

```
    mock_decode, mock_encode, llm_instance
```

```
):
```

```
    mock_encode.return_value = "encoded input with special tokens"
```

```
    mock_decode.return_value = "decoded output with special tokens"
```

```
    result = llm_instance.run("test task with special tokens")
```

```
    mock_encode.assert_called_once()
```

```
    mock_decode.assert_called_once()
```

```
    assert "special tokens" in result
```

```
# Test for correct handling of timeouts
```

```
@patch("swarms.models.huggingface.HuggingfaceLLM.run")
```

```
def test_llm_timeout_handling(mock_run, llm_instance):
```

```
    mock_run.side_effect = TimeoutError
```

```
    with pytest.raises(TimeoutError):
```

```
        llm_instance.run("test task with timeout")
```

```
# Test for response time within a threshold (performance test)
```

```
@patch("swarms.models.huggingface.HuggingfaceLLM.run")
```

```
def test_llm_response_time(mock_run, llm_instance):
```

```
    import time
```

```
    mock_run.return_value = "mocked output"
```

```
    start_time = time.time()
```

```
    llm_instance.run("test task for response time")
```

```
    end_time = time.time()
```

```
    assert (
```

```
        end_time - start_time < 1
```

```
) # Assuming the response should be faster than 1 second
```

```
# Test the logging of a warning for long inputs
```

```
@patch("swarms.models.huggingface.logging.warning")
```

```
def test_llm_long_input_warning(mock_warning, llm_instance):
```

```
    long_input = "x" * 10000 # input longer than the typical limit
```

```

llm_instance.run(long_input)

mock_warning.assert_called_once()


# Test for run method behavior when model raises an exception

@patch(
    "swarms.models.huggingface.HuggingfaceLLM._model.generate",
    side_effect=RuntimeError,
)

def test_llm_run_model_exception(mock_generate, llm_instance):
    with pytest.raises(RuntimeError):
        llm_instance.run("test task when model fails")


# Test the behavior when GPU is forced but not available

@patch("torch.cuda.is_available", return_value=False)

def test_llm_force_gpu_when_unavailable(
    mock_is_available, llm_instance
):
    with pytest.raises(EnvironmentError):
        llm_instance.set_device(
            "cuda"
        ) # Attempt to set CUDA when it's not available


# Test for proper cleanup after model use (releasing resources)

```



```
@patch("swarms.models.huggingface.HuggingfaceLLM._model")

def test_llm_cleanup(mock_model, mock_tokenizer, llm_instance):

    llm_instance.cleanup()

    # Assuming cleanup method is meant to free resources

    mock_model.delete.assert_called_once()

    mock_tokenizer.delete.assert_called_once()
```

# Test model's ability to handle multilingual input

```
@patch("swarms.models.huggingface.HuggingfaceLLM.run")

def test_llm_multilingual_input(mock_run, llm_instance):

    mock_run.return_value = "mocked multilingual output"

    multilingual_input = "Bonjour, ceci est un test multilingue."

    result = llm_instance.run(multilingual_input)

    assert isinstance(

        result, str

    ) # Simple check to ensure output is string type
```

# Test caching mechanism to prevent re-running the same inputs

```
@patch("swarms.models.huggingface.HuggingfaceLLM.run")

def test_llm_caching_mechanism(mock_run, llm_instance):

    input_text = "test caching mechanism"

    mock_run.return_value = "cached output"

    # Run the input twice

    first_run_result = llm_instance.run(input_text)
```

```
second_run_result = llm_instance.run(input_text)

mock_run.assert_called_once() # Should only be called once due to caching

assert first_run_result == second_run_result
```

# These tests are provided as examples. In real-world scenarios, you will need to adapt these tests to the actual logic of your `HuggingfaceLLM` class.

# For instance, "mock\_model.delete.assert\_called\_once()" and similar lines are based on hypothetical methods and behaviors that you need to replace with actual implementations.

# Mock some functions and objects for testing

```
@pytest.fixture
```

```
def mock_huggingface_llm(monkeypatch):
```

```
    # Mock the model and tokenizer creation
```

```
    def mock_init(
```

```
        self,
```

```
        model_id,
```

```
        device="cpu",
```

```
        max_length=500,
```

```
        quantize=False,
```

```
        quantization_config=None,
```

```
        verbose=False,
```

```
        distributed=False,
```

```
        decoding=False,
```

```
        max_workers=5,
```

```
    repitition_penalty=1.3,  
    no_repeat_ngram_size=5,  
    temperature=0.7,  
    top_k=40,  
    top_p=0.8,  
):  
    pass
```

```
# Mock the model loading
```

```
def mock_load_model(self):  
    pass
```

```
# Mock the model generation
```

```
def mock_run(self, task):  
    pass
```

```
monkeypatch.setattr(HuggingfaceLLM, "__init__", mock_init)  
monkeypatch.setattr(HuggingfaceLLM, "load_model", mock_load_model)  
monkeypatch.setattr(HuggingfaceLLM, "run", mock_run)
```

```
# Basic tests for initialization and attribute settings
```

```
def test_init_huggingface_llm():  
    llm = HuggingfaceLLM(  
        model_id="test_model",  
        device="cuda",
```

```
max_length=1000,  
quantize=True,  
quantization_config={"config_key": "config_value"},  
verbose=True,  
distributed=True,  
decoding=True,  
max_workers=3,  
repetition_penalty=1.5,  
no_repeat_ngram_size=4,  
temperature=0.8,  
top_k=50,  
top_p=0.7,  
)
```

```
assert llm.model_id == "test_model"  
assert llm.device == "cuda"  
assert llm.max_length == 1000  
assert llm.quantize is True  
assert llm.quantization_config == {"config_key": "config_value"}  
assert llm.verbose is True  
assert llm.distributed is True  
assert llm.decoding is True  
assert llm.max_workers == 3  
assert llm.repetition_penalty == 1.5  
assert llm.no_repeat_ngram_size == 4  
assert llm.temperature == 0.8
```

```
assert llm.top_k == 50
```

```
assert llm.top_p == 0.7
```

```
# Test loading the model
```

```
def test_load_model(mock_huggingface_llm):
```

```
    llm = HuggingfaceLLM(model_id="test_model")
```

```
    llm.load_model()
```

```
# Test running the model
```

```
def test_run(mock_huggingface_llm):
```

```
    llm = HuggingfaceLLM(model_id="test_model")
```

```
    llm.run("Test prompt")
```

```
# Test for setting max_length
```

```
def test_llm_set_max_length(llm_instance):
```

```
    new_max_length = 1000
```

```
    llm_instance.set_max_length(new_max_length)
```

```
    assert llm_instance.max_length == new_max_length
```

```
# Test for setting verbose
```

```
def test_llm_set_verbose(llm_instance):
```

```
    llm_instance.set_verbose(True)
```

```
assert llm_instance.verbose is True
```

```
# Test for setting distributed
```

```
def test_llm_set_distributed(llm_instance):
```

```
    llm_instance.set_distributed(True)
```

```
    assert llm_instance.distributed is True
```

```
# Test for setting decoding
```

```
def test_llm_set_decoding(llm_instance):
```

```
    llm_instance.set_decoding(True)
```

```
    assert llm_instance.decoding is True
```

```
# Test for setting max_workers
```

```
def test_llm_set_max_workers(llm_instance):
```

```
    new_max_workers = 10
```

```
    llm_instance.set_max_workers(new_max_workers)
```

```
    assert llm_instance.max_workers == new_max_workers
```

```
# Test for setting repitition_penalty
```

```
def test_llm_set_repitition_penalty(llm_instance):
```

```
    new_repitition_penalty = 1.5
```

```
    llm_instance.set_repitition_penalty(new_repitition_penalty)
```

```
assert llm_instance.repetition_penalty == new_repetition_penalty
```

```
# Test for setting no_repeat_ngram_size
```

```
def test_llm_set_no_repeat_ngram_size(llm_instance):
```

```
    new_no_repeat_ngram_size = 6
```

```
    llm_instance.set_no_repeat_ngram_size(new_no_repeat_ngram_size)
```

```
    assert (
```

```
        llm_instance.no_repeat_ngram_size == new_no_repeat_ngram_size
```

```
)
```

```
# Test for setting temperature
```

```
def test_llm_set_temperature(llm_instance):
```

```
    new_temperature = 0.8
```

```
    llm_instance.set_temperature(new_temperature)
```

```
    assert llm_instance.temperature == new_temperature
```

```
# Test for setting top_k
```

```
def test_llm_set_top_k(llm_instance):
```

```
    new_top_k = 50
```

```
    llm_instance.set_top_k(new_top_k)
```

```
    assert llm_instance.top_k == new_top_k
```

# Test for setting top\_p

```
def test_llm_set_top_p(llm_instance):
```

```
    new_top_p = 0.9
```

```
    llm_instance.set_top_p(new_top_p)
```

```
    assert llm_instance.top_p == new_top_p
```

# Test for setting quantize

```
def test_llm_set_quantize(llm_instance):
```

```
    llm_instance.set_quantize(True)
```

```
    assert llm_instance.quantize is True
```

# Test for setting quantization\_config

```
def test_llm_set_quantization_config(llm_instance):
```

```
    new_quantization_config = {
```

```
        "load_in_4bit": False,
```

```
        "bnb_4bit_use_double_quant": False,
```

```
        "bnb_4bit_quant_type": "nf4",
```

```
        "bnb_4bit_compute_dtype": torch.bfloat16,
```

```
    }
```

```
    llm_instance.set_quantization_config(new_quantization_config)
```

```
    assert llm_instance.quantization_config == new_quantization_config
```

# Test for setting model\_id



```
def test_llm_set_model_id(llm_instance):

    new_model_id = "EleutherAI/gpt-neo-2.7B"

    llm_instance.set_model_id(new_model_id)

    assert llm_instance.model_id == new_model_id


# Test for setting model

@patch(

    "swarms.models.huggingface.AutoModelForCausalLM.from_pretrained"

)

def test_llm_set_model(mock_model, llm_instance):

    mock_model.return_value = "mocked model"

    llm_instance.set_model(mock_model)

    assert llm_instance.model == "mocked model"


# Test for setting tokenizer

@patch("swarms.models.huggingface.AutoTokenizer.from_pretrained")

def test_llm_set_tokenizer(mock_tokenizer, llm_instance):

    mock_tokenizer.return_value = "mocked tokenizer"

    llm_instance.set_tokenizer(mock_tokenizer)

    assert llm_instance.tokenizer == "mocked tokenizer"


# Test for setting logger

def test_llm_set_logger(llm_instance):
```

```
new_logger = logging.getLogger("test_logger")

llm_instance.set_logger(new_logger)

assert llm_instance.logger == new_logger
```

# Test for saving model

```
@patch("torch.save")
```

```
def test_llm_save_model(mock_save, llm_instance):
```

```
    llm_instance.save_model("path/to/save")
```

```
    mock_save.assert_called_once()
```

# Test for print\_dashboard

```
@patch("builtins.print")
```

```
def test_llm_print_dashboard(mock_print, llm_instance):
```

```
    llm_instance.print_dashboard("test task")
```

```
    mock_print.assert_called()
```

# Test for \_\_call\_\_ method

```
@patch("swarms.models.huggingface.HuggingfaceLLM.run")
```

```
def test_llm_call(mock_run, llm_instance):
```

```
    mock_run.return_value = "mocked output"
```

```
    result = llm_instance("test task")
```

```
    assert result == "mocked output"
```